

Deep Learning

Lecture 5: Generative Adversarial, Variational and Implicit Models

Amir Atapour-Abarghouei

amir.atapour-abarghouei@durham.ac.uk

Durham University





Lecture Overview

- Generative models
 - Unsupervised training
 - PixelRNN and PixelCNN

Generative adversarial networks

- Definition
- Common Issues
- Lipschitz Continuity
- Spectral Normalisation
- Conditional GANs
- Other Advanced Variants

Variational autoencoders

- Autoencoders
- VAEs
- Reparameterisation trick
- ELBO

Implicit networks


- Definition
- SIRENs
- NeRFs



Supervised vs Unsupervised Training

- Supervised Learning:

- Data x – Labels y
- Learn mapping function $x \rightarrow y$
- e.g. classification, regression, segmentation, object detection, machine translation, sentiment analysis (and everything else we have talked about so far).



- Unsupervised Learning:

- Data x – **No** Labels
- Learn underlying structures and patterns in the data.
- e.g. clustering, dimensionality reduction, feature learning, density estimation.
- **Generative models** are **unsupervised** as the data used has **no labels**.





Generative Models

Definition:

Given a set of data, capture the probability distribution representing this data and generate new data samples from this distribution.



Training Data $\sim P_{\text{data}}(x)$

<https://github.com/NVlabs/ffhq-dataset>



Generated Sample $\sim P_{\text{model}}(x)$

<https://arxiv.org/pdf/1912.04958.pdf>


The objective is to learn $P_{\text{model}}(x)$ so it is as close as possible to $P_{\text{data}}(x)$.



- Explicit tractable density function.
- We generate pixels one by one starting from a corner.
- Dependency on previous pixels is modelled via an RNN.
- During inference, for generation, we initialise the value of the first pixel and the image is generated.

Good results

but RNNs are sequential and thus **inefficient!**



... continue until the entire image is generated.

PixelCNN


[van den Oord et al., 2016]



- Explicit tractable density function.
- We generate pixels one by one starting from a corner.
- Dependency on **pixels** is modelled via a CNN over context patch.

CNN is parallelisable and thus **efficient!**

During inference, generation is sequential and still very slow...!




... continue until the entire image is generated.

PixelRNN / PixelCNN



[van den Oord et al., 2016]

PixelRNN



PixelCNN






Generative Adversarial Networks

Definition:

A generative adversarial network (GAN) is a non-cooperative zero-sum game where two networks compete against each other


[Goodfellow et al., 2014].

One network $G(z)$ generates new samples, whereas D estimates the probability the sample was from the training data rather than G :



$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

Generative Adversarial Networks



Generator

1. →

I see!!



2. →

Hmm!!



3. →

Aha!!



4. →

Success!



Real?

Fake!

Fake!

Fake?




Discriminator

Generative Adversarial Networks

Training




We always start from z .
(Random Noise)



Note that when the Discriminator is being trained, the weights of the Generator are not updated.

First, we need to train the **Discriminator**:

1. Use a random noise vector with the **Generator** to produce a fake image.
2. Use the fake image and a real image to only train the **Discriminator**.




Generative Adversarial Networks

Training



Again, we sample a z .



Next, we need to train the **Generator**:

1. Use a random noise vector with the **Generator** to produce a fake image.
2. Pass the fake image through the **Discriminator**, but with a **Real** label.
3. We then use the gradients from the **Discriminator** to train the **Generator**.

Note that we are trying to **fool** the Discriminator, which is why the **Real** label is used for the fake image.

Generative Adversarial Networks

Objective



- The objective function:
 - Generator and Discriminator are trained jointly in minimax game:

$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

Discriminator output for real data x Discriminator output for fake generated data $G(z)$

- Discriminator (θ_d) wants to **maximise** the objective so $D(x)$ is close to 1 (real) and $D(G(z))$ is close to 0 (fake).
- Generator (θ_g) wants to **minimise** the objective so $D(G(z))$ is close to 1 (the discriminator will think the fake output is real).

Let's Look at Some Code!



Simple GAN Example

Code on GitHub : https://github.com/atapour/dl-pytorch/blob/main/Simple_GAN_Example/Simple_GAN_Example.ipynb

Code on Colab: https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Simple_GAN_Example/Simple_GAN_Example.ipynb



Issues:

- **Non-Convergence:** Model is unstable so parameters oscillate and never converge.
- **Diminishing Gradients:** The Discriminator reaches its optimal state too soon and its small gradients cannot train the Generator.
- **Hyperparameter Sensitivity:** The model is too sensitive to any changes in hyperparameters (e.g., learning rate).
- **Mode Collapse:** The Generator collapses and produces a limited variety of samples.

Generative Adversarial Networks

Improvements



Definition: Lipschitz function

A function f is Lipschitz continuous if it is bounded by how fast it can change. Specifically if there exists a positive real constant k where:

$$|f(x) - f(y)| \leq k|x - y|,$$

for all y sufficiently near x . For example, any function with a bounded first derivative is a Lipschitz function.

The slope of any secant line to f is between $-k$ and k

Lower Lipschitz constant makes the function smooth and easy to optimise.

If the inputs x and y are close/similar, we want model outputs, $f(x)$ and $f(y)$ to also be similar.

Generative Adversarial Networks

Improvements



- Wasserstein GAN [Arjovsky et al., 2017]
 - Limiting the gradients of the discriminator will make it a smoother function.
 - Done by clipping the gradients.
- Improved Wasserstein GAN [Gulrajani et al., 2017]
 - Add a penalty term to penalise the gradients of the discriminator directly.
 - Requires second order differentiation, which PyTorch supports, but it is slow.

Generative Adversarial Networks

Improvements




- There is another solution using Normalisation!

First, what is normalisation? – BatchNorm [Ioffe and Szegedy, 2015]

- All layers are updated in the backward pass.
- Every layer assumes other layers are fixed.

The weights of a layer can be updated based on the expectation that the prior layer produces values with a given distribution. However, this distribution is likely to change once that layer is updated.



Batch Normalisation: Layer outputs are rescaled, so they have a mean of *zero* and a standard deviation of *one*:

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

- Weight Norm
- Layer Norm
- Instance Norm
- Group Norm
- Switchable Norm
- Etc.

Problem: Internal Covariate Shift

(change in the distribution of input during training)

Generative Adversarial Networks

Improvements



- Spectral Normalisation [Miyato et al., 2018]

Definition: spectral normalisation

The matrix (spectral) norm defines how much a matrix can stretch a vector x :

$$\| A \| = \max_{x \neq 0} \frac{\| Ax \|}{\| x \|}$$

Spectral norm normalises the weights for each layer using the spectral norm $\sigma(W)$ such that the Lipschitz constant for every layer and the whole network is 1.

$$\hat{W}_{SN} = W / \sigma(W)$$

$$\sigma(\hat{W}_{SN}(W)) = 1$$

$$\| f \|_{Lip} = 1$$

Pseudocode: 1-Lipschitz discriminator

```
class Discriminator(nn.Module):
    def __init__(self, f=64):
        super().__init__()
        self.discriminate = nn.Sequential(
            spectral_norm(Conv2d(1, f, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f, f*2, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f*2, f*4, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f*4, f*8, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f*8, 1, 3, 1, 1)),
            nn.Sigmoid()
        )
```

Generative Adversarial Networks

Improvements



- **DCGAN:** The entire architecture is fully convolutional.
- **Alternative Objective:** Wasserstein GAN - LSGAN,
- **Two Timescale Update:** The Generator is updated more slowly than the Discriminator (SAGAN).
Mean squared loss
- **Stacking GANs:** Multiple GANs are placed consecutively, and each GAN solves an easier version of the problem (FashionGAN).
- **Progressive Growing GAN:** Higher quality and larger outputs by incrementally increasing the size of the model during training.


Generative Adversarial Networks

Conditional
GANs



Definition: Conditional GANs

GANs can be conditioned, for instance with labels y , if available, [Mirza et al., 2014] by feeding the label information into both the generator and the discriminator:



$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x|y) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z, y|y)))]$$



Let's Look at Some Code!

Conditional GAN Example

Code on GitHub : https://github.com/atapour/dl-pytorch/blob/main/Conditional_GAN_Example/Conditional_GAN_Example.ipynb

Code on Colab: https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Conditional_GAN_Example/Conditional_GAN_Example.ipynb

Generative Adversarial Networks

Advanced
Variants



- Auxiliary Classifier GAN (ACGAN)
- InfoGAN
- Least Square GAN (LSGAN)
- Adversarial Autoencoder
- Boundary Equilibrium GAN (BEGAN)
- Energy-Based GAN
- StyleGAN
- etc. etc. etc. ...


Popular Applications

Unpaired Translation (CycleGAN)



Definition: CycleGAN [Zhu et al., 2017]

An adversarial architecture for unpaired image translation. It has twin generators with skip connections and two discriminators, which translate between the domains, alongside a cycle consistency loss (L1 distance) to ensure the mapping recovers the original image.






Popular Applications

Super-resolution

Definition: Super-resolution

The objective is to map a single low-resolution input to a distribution of high-resolution outputs. PULSE [Menon et al., 2020] projects points in the search of the latent space of StyleGAN (a large conditional GAN) onto a hypersphere, which ensures probable outputs in the high-dimensional latent space.



Popular Applications

Anomaly Detection




Definition: Anomaly Detection

The objective is to find “anomalies” in the data. Unsupervised anomaly detectors [1,2] learn a normal distribution over (healthy) observations. Then, when they observe something not observed in training (unhealthy/dangerous), they fail to reconstruct - detecting it as an anomaly. Region-based anomaly detectors [3] learn a distribution over inpainted (erased) regions.

[1] Akçay, **Atapour-Abarghouei**, and Breckon. “GANomaly: Semi-Supervised Anomaly Detection via Adversarial Training”, ACCV, 2018.

[2] Akçay, **Atapour-Abarghouei**, and Breckon. “Skip-ganomaly: Skip connected and adversarially trained encoder-decoder anomaly detection”, IJCNN, 2019.

[3] Nguyen, Feldman, Bethapudi, Jennings, and Willcocks. “Unsupervised Region-based Anomaly Detection in Brain MRI with Adversarial Image Inpainting”, arXiv:2010.01942.






Autoencoders

- Not generative models.
- Designed to obtain latent variables.
- Latent variables inferred by an autoencoder are a lower dimensional feature representation from unlabeled data.

variables that are not directly observed but are rather inferred (through a model)






Autoencoders

Definition: Autoencoder

An autoencoder is a feedforward neural network that reconstructs its inputs, thus learning an identity:

$$\mathcal{L}_{AE} = E_{x \sim p_{\text{data}}} [\mathcal{L}(x, D(E(x)))]$$

where \mathcal{L} is a loss function. The encoder function $E : R_n \rightarrow R_m$ compresses the dimensionality of the data $n \ll m$ to a latent encoding $z = E(x)$, which is then recovered by the decoder $D : R_m \rightarrow R_n$, where $\hat{x} = D(z)$.






Autoencoders


Properties: Autoencoder


- Very good at representation learning.
- Are being used for data compression.

- Denoising Autoencoders

 Reconstruct image from a noisy input.

- Inpainting Autoencoders

 Reconstruct image from a input with missing regions.



Let's Look at Some Code!



Example of a Simple Autoencoder

Code on GitHub: https://github.com/atapour/dl-pytorch/blob/main/AutoEncoder_Example/AutoEncoder_Example.ipynb


Code on Colab: https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/AutoEncoder_Example/AutoEncoder_Example.ipynb

Variational Autoencoders (VAE)

[Kingma and Welling, 2013]



- Explicit approximate variational density function.
- ***Generative Models*** with a probabilistic spin on autoencoders.
- Instead of a fixed latent vector (as in an autoencoder), **VAEs learn a probability distribution** in the latent space that facilitates sampling new data for generation.



Variational Autoencoders (VAE)

[Kingma and Welling, 2013]





- The encoder learns:
 - vector μ (**means**)
 - vector σ (**standard deviations**)
- The loss function will include a reconstruction loss, just like the vanilla autoencoder, but there is more



the latent variables in z are sampled

describe

The Encoder computes $p_\phi(z|x)$. The Decoder computes $q_\theta(x|z)$.



$$\mathcal{L}(x, \hat{x}) = \| x, \hat{x} \|^2$$

Variational Autoencoders (VAE)

[Kingma and Welling, 2013]




- The loss function consists of two terms:

- Reconstruction Loss (L_2 loss)
 - KL Divergence


$$\mathcal{L}(\phi, \theta; x, z) = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{\text{KL}}(q_\phi(z|x) \parallel p(z))$$

Expectation: we are dealing with random variables

Reminder:
summation or
integration of all
possible values from
a random variable



The Encoder
 $q_\phi(z|x)$.



Kullback-Leibler divergence: measure of how one probability distribution is different from another



The Decoder
 $p_\theta(x|z)$.

Variational Autoencoders (VAE)

[Kingma and
Welling, 2013]



We want the probability distribution to be smooth and not to overfit or memorise any of the latent variables within vector z .

$$\mathcal{L}(\phi, \theta; x, z) = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - \boxed{\text{D}_{\text{KL}}(q_\phi(z|x) \parallel p(z))}$$

- By reducing the distance between the probability distribution and a well-behaved prior distribution (such as a *Normal Gaussian*), we:
 - encourage the latent variables to be evenly distributed around the centre of the latent space
 - and deter the model from clustering points in specific regions and cheat by memorising the data.

Variational Autoencoders (VAE)

[Kingma and Welling, 2013]




Definition: Kullback–Leibler divergence

The Kullback–Leibler divergence (also called relative entropy) measures the difference between distributions and is asymmetric and non-negative:

$$D_{\text{KL}}(p \parallel q) = \int p(x) \log \left(\frac{p(x)}{q(x)} \right) dx$$

Where $D_{\text{KL}}(p \parallel p) = 0$. Practically, the KL divergence is sensitive at the tails of the distribution.



Variational Autoencoders (VAE)

[Kingma and Welling, 2013]




Let's visually inspect how regularisation with the KL term helps.

- We can analyse the effect of the *KL* term by weighting the loss and experimenting with different values:

$$\text{loss} = \text{reconstruction_loss} + c * \text{kl_loss}$$

- We will consider an autoencoder with a latent space of 2 dimensions trained on the MNIST dataset.

For reference, let's look at points sampled from the *standard Gaussian* in two dimensions.




Variational Autoencoders (VAE)

[Kingma and Welling, 2013]




We can see how the points begin to cluster with *no* or *a small k1_loss*.


C = 0.0




C = 0.02




C = 0.1




C = 0.5



C = 1.0



C = 2.0



Variational Autoencoders (VAE)

[Kingma and Welling, 2013]




Issue to overcome: backpropagation through a sampling layer

Solution: **Reparameterisation Trick**

$$z = \mu + \sigma \odot \varepsilon \quad \text{where } \varepsilon \sim \mathcal{N}(0,1)$$

z will be a sum of a fixed vector μ and a fixed vector σ scaled by random constants ε sampled from the prior distribution (standard Gaussian).




Variational Autoencoders (VAE)

[Kingma and Welling, 2013]



- Reparameterisation Trick



Variational Autoencoders (VAE)

[Kingma and Welling, 2013]



Definition: Evidence Lower Bound (ELBO)


VAEs have three components:

1. The decoder $p_\theta(x|z)$
2. The approximate posterior (encoder) $q_\phi(z|x)$
3. The prior distribution $p_\theta(z)$

They are trained with the reparameterisation trick to maximise the evidence lower bound (ELBO):

$$\log p_\theta(x) \geq \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x) \parallel p(z))$$

There are more advanced and state-of-the-art method, such stacking VAEs hierarchically [Child et al., 2021].



Variational Autoencoders (VAE)

[Kingma and Welling, 2013]




- **Results**

The latent variables are forced to be independent of each other.



Different dimensions
of z encode
interpretable factors of
variation.



Let's Look at Some Code!



Example of a Variational Autoencoder

Code on GitHub: https://github.com/atapour/dl-pytorch/blob/main/VAE_Example/VAE_Example.ipynb

Code on Colab: https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/VAE_Example/VAE_Example.ipynb



Implicit Models

Definition


Definition: implicit representations

Consider data $\Phi: \mathbb{R}^m \rightarrow \mathbb{R}^n$, like a single image, as a function of coordinates $\mathbf{c} \in \mathbb{R}^m$. The aim is to learn a neural approximation of Φ that satisfies an implicit equation:

$$R(\mathbf{c}, \Phi, \nabla_\Phi, \nabla_\Phi^2, \dots) = 0, \Phi: \mathbf{c} \rightarrow \Phi(\mathbf{c})$$


This has applications in 3D modelling, image, video, audio representation.

Example: Implicit Network







Implicit Models



coordinates - (x, y)



pixel value - $\{R, G, B\}$



We can have a neural network approximate the image function.

Implicit Representation Networks SIREN




Definition: SIREN

SInusoidal REpresentation Networks (SIREN) are a simple implicit representation network, but use **sin** (with clever initialisation) as their choice of non-linearity [Sitzmann et al., 2020].

sin is periodic, so it allows to capture patterns over all of the coordinate space (translation invariant, like convolutions).

Example: SIREN



<https://www.vincentsitzmann.com/siren>

Implicit Representation Networks




NeRF

Definition: NeRF

Neural Radiance Fields (NeRF) are similar to SIRENs, but instead of representing an image, they represent a single 3D scene [Mildenhall et al., 2020].

They map from pixel positions (x, y, z) and a viewing direction (θ, ϕ) to a colour and density value σ integrated via a ray on F_θ .

Example: NeRF



<https://www.matthewtancik.com/nerf>

There are other implicit models:

<https://cwkx.github.io/data/GON/>



What we learned today!

- Generative models
 - Unsupervised training
 - PixelRNN and PixelCNN

Generative adversarial networks

- Definition
- Common Issues
- Lipschits Continuity
- Spectral Normalisation
- Conditional GANs
- Other Advanced Variants

Variational autoencoders

- Autoencoders
- VAEs
- Reparameterisation trick
- ELBO

Implicit networks

- Definition
- SIRENs
- NeRFs