

# Deep Learning

## Lecture 4: Generative *Adversarial* Models

---

Amir Atapour-Abarghouei

*amir.atapour-abarghouei@durham.ac.uk*

Durham University





# Lecture Overview

## 1 Generative models

- Unsupervised training
- PixelRNN and PixelCNN

## 2 Generative adversarial networks

- Definition
- Common Issues
- Lipschits Continuity
- Spectral Normalisation
- Conditional GANs
- Other Advanced Variants

## 3 Popular applications

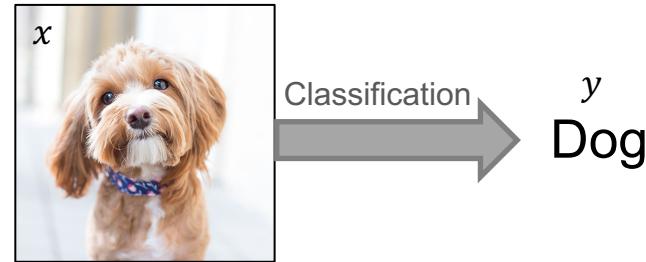
- Unpaired Translation
- Super Resolution
- Anomaly Detection



# Supervised vs Unsupervised Training

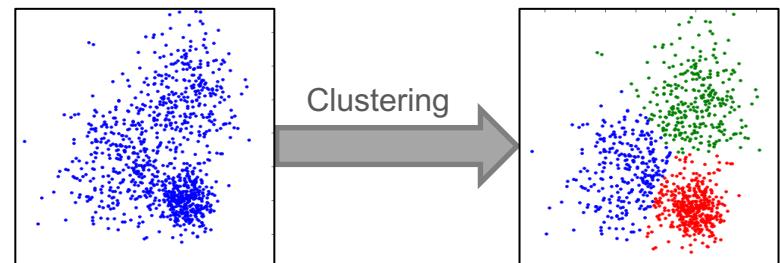
## • Supervised Learning:

- Data  $x$  – Labels  $y$
- Learn mapping function  $x \rightarrow y$
- e.g. classification, regression, segmentation, object detection, machine translation, sentiment analysis (and everything else we have talked about so far).



## • Unsupervised Learning:

- Data  $x$  – **No** Labels
- Learn underlying structures and patterns in the data.
- e.g. clustering, dimensionality reduction, feature learning, density estimation.
- **Generative models** are **unsupervised** as the data used has **no labels**.





# Generative Models

## Definition:

Given a set of data, capture the probability distribution representing this data and generate new data samples from this distribution.



Training Data  $\sim P_{\text{data}}(x)$

<https://github.com/NVlabs/ffhq-dataset>



Generated Sample  $\sim P_{\text{model}}(x)$

<https://arxiv.org/pdf/1912.04958.pdf>

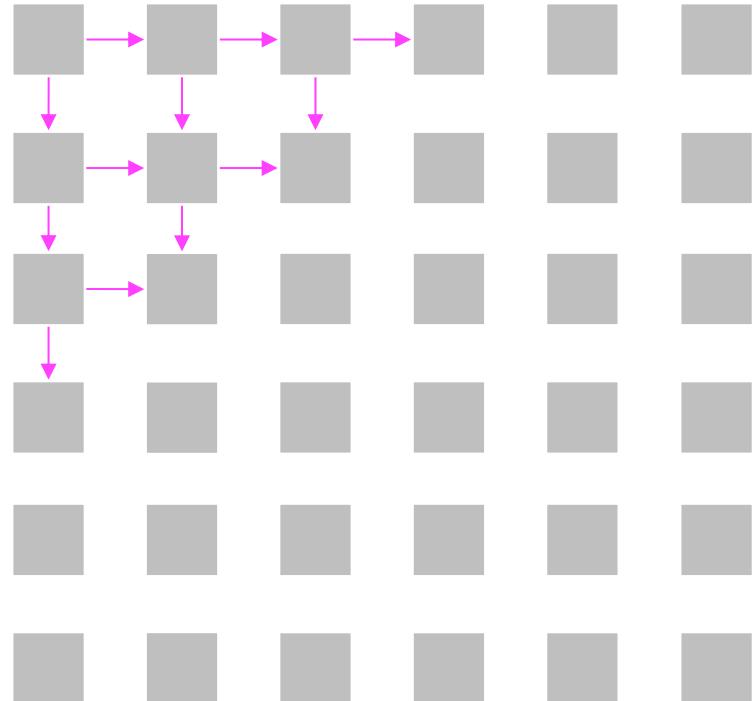
The objective is to learn  $P_{\text{model}}(x)$  so it is as close as possible to  $P_{\text{data}}(x)$ .



- Explicit tractable density function.
- We generate pixels one by one starting from a corner.
- Dependency on previous pixels is modelled via an RNN.
- During inference, for generation, we initialise the value of the first pixel and the image is generated.

Good results

but RNNs are sequential and thus **inefficient!**



*... continue until the entire image is generated.*

# PixelCNN

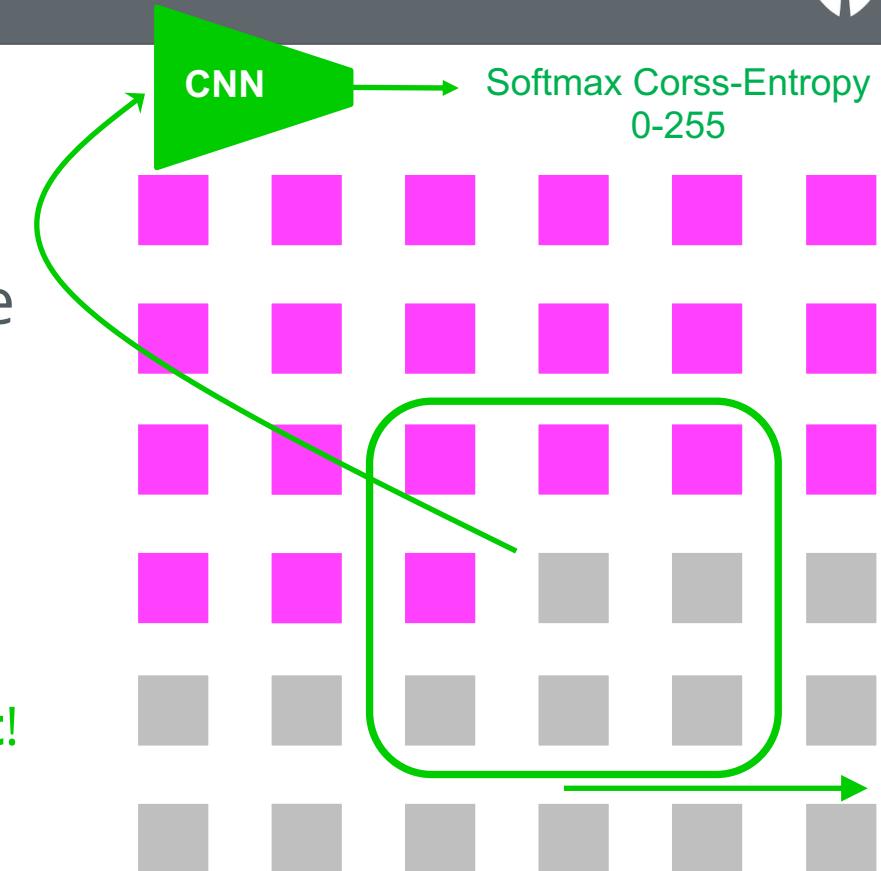
[van den Oord et al., 2016]



- Explicit tractable density function.
- We generate pixels one by one starting from a corner.
- Dependency on **pixels** is modelled via a CNN over context patch.

CNN is parallelisable and thus **efficient!**

During inference, generation is sequential and still very slow...!



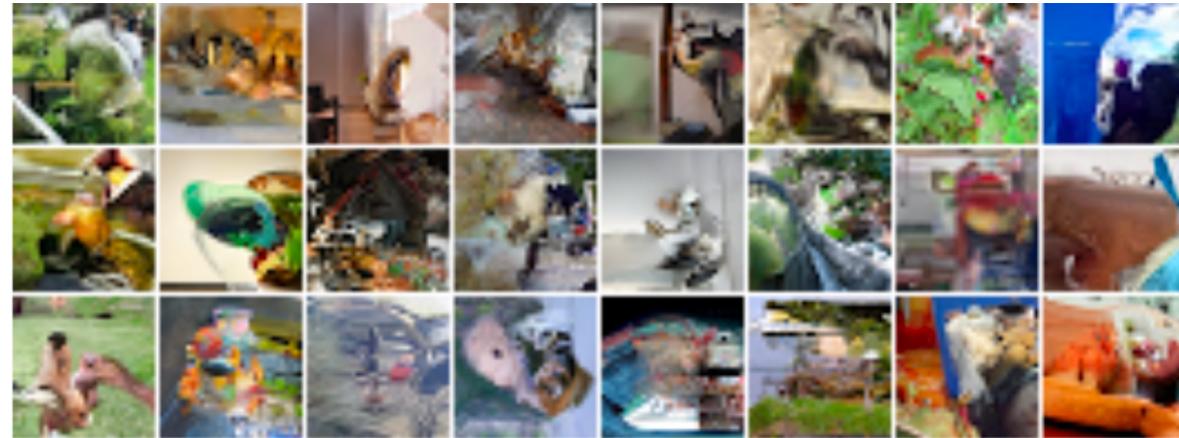
*... continue until the entire image is generated.*

# PixelRNN / PixelCNN



[van den Oord et al., 2016]

PixelRNN



PixelCNN





# Generative Adversarial Networks

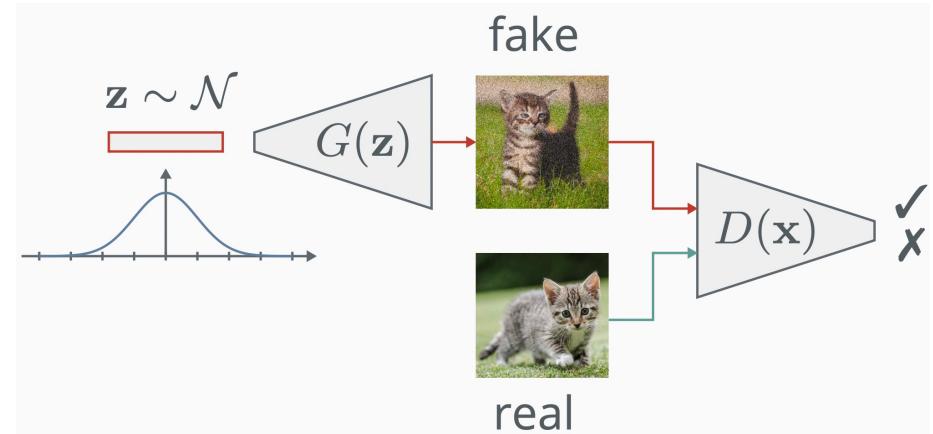
## Definition:

A generative adversarial network (GAN) is a non-cooperative zero-sum game where two networks compete against each other

[Goodfellow et al., 2014].

One network  $G(z)$  generates new samples, whereas  $D$  estimates the probability the sample was from the training data rather than  $G$ :

$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$





# Generative Adversarial Networks



Generator

1. →

I see!!



Fake!

2. →

Hmm!!



Fake!

3. →

Aha!!



Fake?

4. →

Success!



Real?



Discriminator

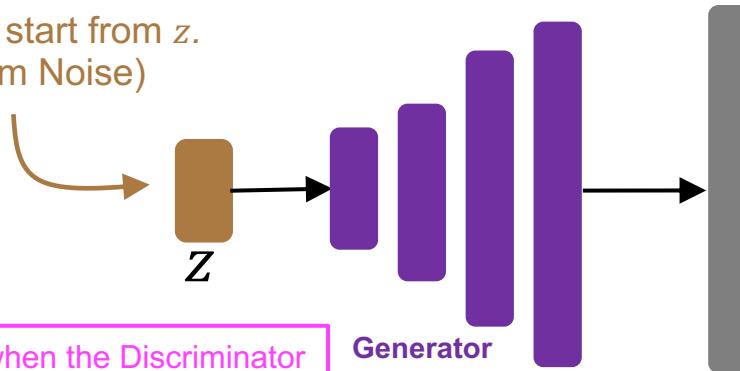


# Generative Adversarial Networks

Training



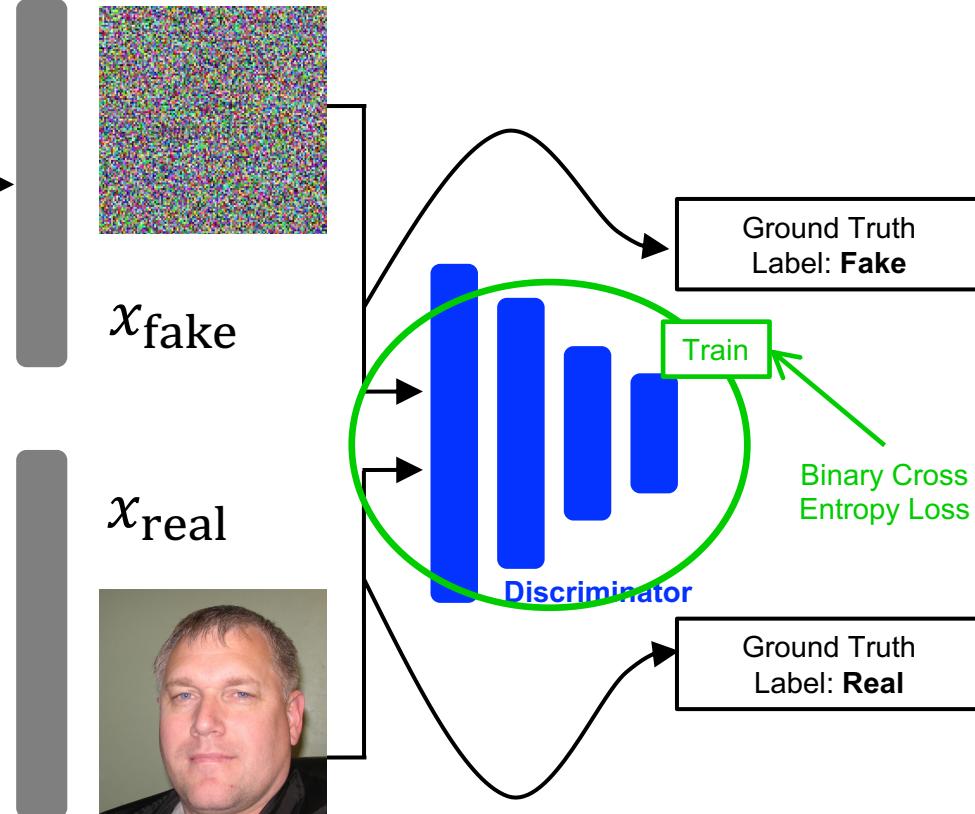
We always start from  $z$ .  
(Random Noise)



Note that when the Discriminator is being trained, the weights of the Generator are not updated.

First, we need to train the **Discriminator**:

1. Use a random noise vector with the **Generator** to produce a fake image.
2. Use the fake image and a real image to only train the **Discriminator**.

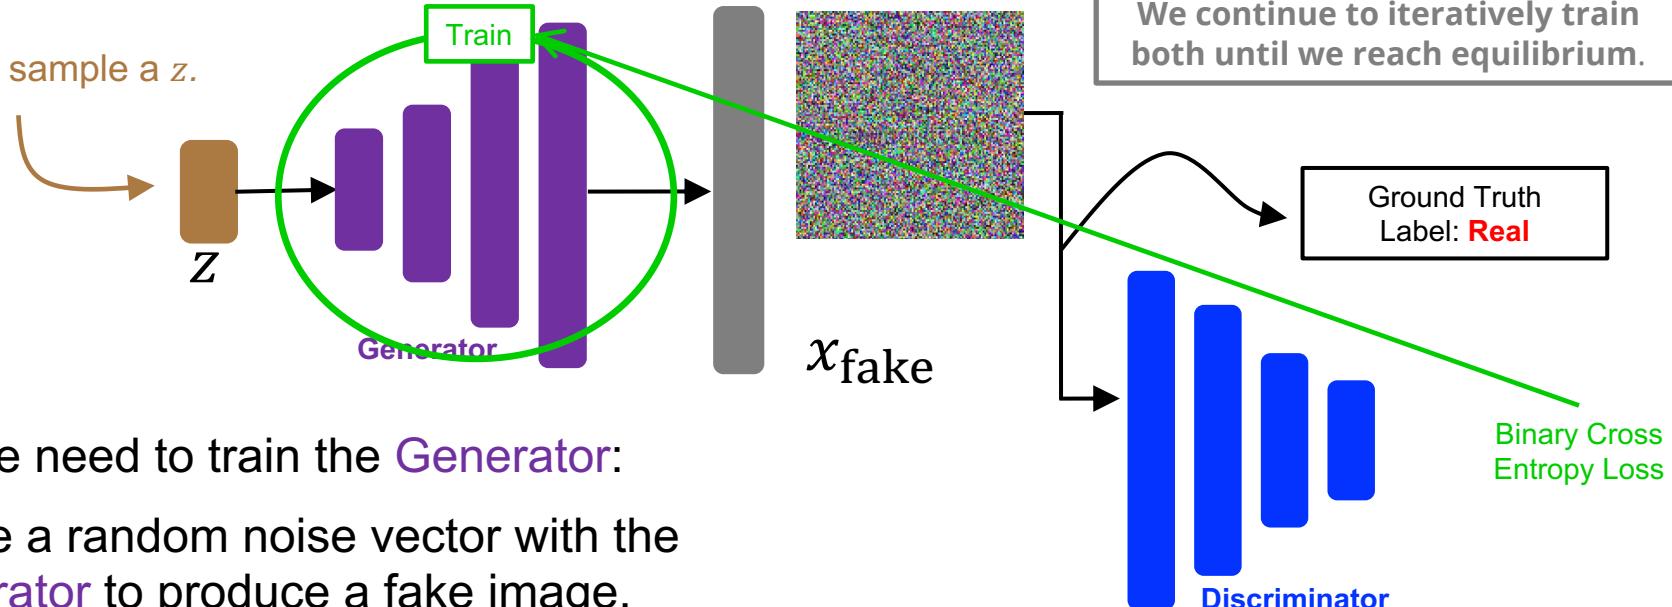


# Generative Adversarial Networks

Training



Again, we sample a  $z$ .



Next, we need to train the **Generator**:

1. Use a random noise vector with the **Generator** to produce a fake image.
2. Pass the fake image through the **Discriminator**, but with a **Real** label.
3. We then use the gradients from the **Discriminator** to train the **Generator**.

Note that we are trying to **fool** the Discriminator, which is why the **Real** label is used for the fake image.

# Generative Adversarial Networks

Objective



- The objective function:
  - Generator and Discriminator are trained jointly in minimax game:

$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$

Discriminator output for real data  $x$       Discriminator output for fake generated data  $G(z)$

- Discriminator ( $\theta_d$ ) wants to **maximise** the objective so  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake).
- Generator ( $\theta_g$ ) wants to **minimise** the objective so  $D(G(z))$  is close to 1 (the discriminator will think the fake output is real).

# Let's Look at Some Code!



## Simple GAN Example

Code on GitHub : [https://github.com/atapour/dl-pytorch/blob/main/Simple\\_GAN\\_Example/Simple\\_GAN\\_Example.ipynb](https://github.com/atapour/dl-pytorch/blob/main/Simple_GAN_Example/Simple_GAN_Example.ipynb)

Code on Colab: [https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Simple\\_GAN\\_Example/Simple\\_GAN\\_Example.ipynb](https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Simple_GAN_Example/Simple_GAN_Example.ipynb)



## Issues:

- **Non-Convergence:** Model is unstable so parameters oscillate and never converge.
- **Diminishing Gradients:** The Discriminator reaches its optimal state too soon and its small gradients cannot train the Generator.
- **Hyperparameter Sensitivity:** The model is too sensitive to any changes in hyperparameters (e.g., learning rate).
- **Mode Collapse:** The Generator collapses and produces a limited variety of samples.

# Generative Adversarial Networks

Improvements



## Definition: Lipschitz function

A function  $f$  is Lipschitz continuous if it is bounded by how fast it can change. Specifically if there exists a positive real constant  $k$  where:

$$|f(x) - f(y)| \leq k|x - y|,$$

for all  $y$  sufficiently near  $x$ . For example, any function with a bounded first derivative is a Lipschitz function.

The slope of any secant line to  $f$  is between  $-k$  and  $k$

Lower Lipschitz constant makes the function smooth and easy to optimise.

If the inputs  $x$  and  $y$  are close/similar, we want model outputs,  $f(x)$  and  $f(y)$  to also be similar.

# Generative Adversarial Networks

Improvements



- Wasserstein GAN [Arjovsky et al., 2017]
  - Limiting the gradients of the discriminator will make it a smoother function.
  - Done by clipping the gradients.
- Improved Wasserstein GAN [Gulrajani et al., 2017]
  - Add a penalty term to penalise the gradients of the discriminator directly.
  - Requires second order differentiation, which PyTorch supports, but it is slow.

# Generative Adversarial Networks

Improvements



- There is another solution using Normalisation!

First, what is normalisation? – BatchNorm [Ioffe and Szegedy, 2015]

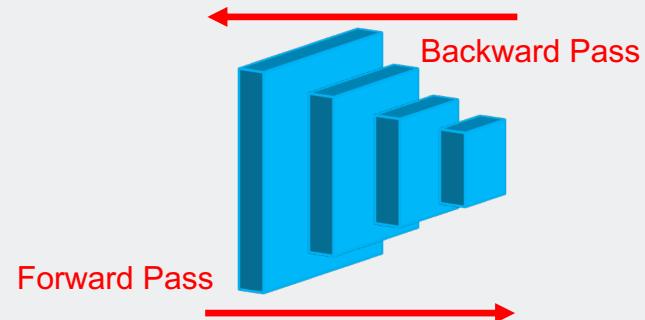
- All layers are updated in the backward pass.
- Every layer assumes other layers are fixed.

The weights of a layer can be updated based on the expectation that the prior layer produces values with a given distribution. However, this distribution is likely to change once that layer is updated.

**Batch Normalisation:** Layer outputs are rescaled, so they have a mean of zero and a standard deviation of *one*:

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

- Weight Norm
- Layer Norm
- Instance Norm
- Group Norm
- Switchable Norm
- Etc.



**Problem: Internal Covariate Shift**

(change in the distribution of input during training)

# Generative Adversarial Networks

Improvements



- Spectral Normalisation [Miyato et al., 2018]

## Definition: spectral normalisation

The matrix (spectral) norm defines how much a matrix can stretch a vector  $x$ :

$$\| A \| = \max_{x \neq 0} \frac{\| Ax \|}{x}$$

Spectral norm normalises the weights for each layer using the spectral norm  $\sigma(W)$  such that the Lipschitz constant for every layer and the whole network is 1.

$$\hat{W}_{\text{SN}} = W / \sigma(W)$$

$$\sigma(\hat{W}_{\text{SN}}(W)) = 1$$

$$\| f \|_{\text{Lip}} = 1$$

## Pseudocode: 1-Lipschitz discriminator

```
class Discriminator(nn.Module):
    def __init__(self, f=64):
        super().__init__()
        self.discriminate = nn.Sequential(
            spectral_norm(Conv2d(1, f, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f, f*2, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f*2, f*4, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f*4, f*8, 3, 1, 1)),
            nn.LeakyReLU(0.1, inplace=True),
            nn.MaxPool2d(kernel_size=(2,2)),
            spectral_norm(Conv2d(f*8, 1, 3, 1, 1)),
            nn.Sigmoid()
        )
```

# Generative Adversarial Networks

Improvements



- **DCGAN:** The entire architecture is fully convolutional.
- **Alternative Objective:** Wasserstein GAN - LSGAN,
- **Two Timescale Update:** The Generator is updated more slowly than the Discriminator (SAGAN).  
Mean squared loss
- **Stacking GANs:** Multiple GANs are placed consecutively, and each GAN solves an easier version of the problem (FashionGAN).
- **Progressive Growing GAN:** Higher quality and larger outputs by incrementally increasing the size of the model during training.

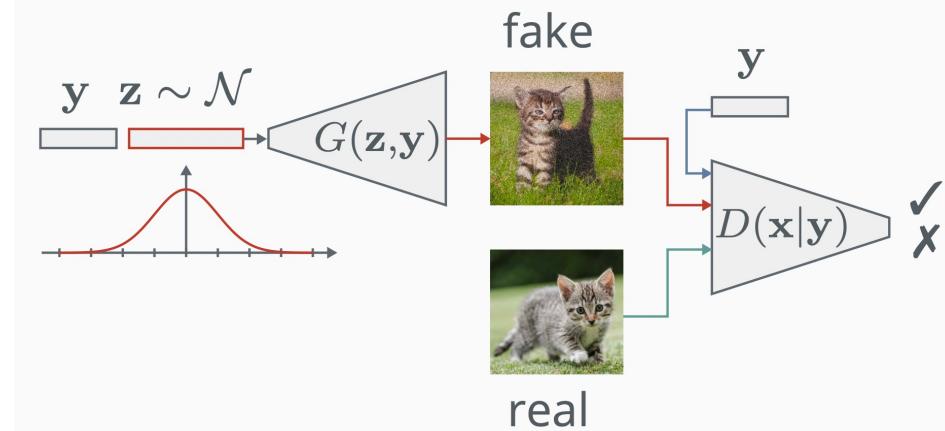
# Generative Adversarial Networks

Conditional  
GANs



## Definition: Conditional GANs

GANs can be conditioned, for instance with labels  $y$ , if available, [Mirza et al., 2014] by feeding the label information into both the generator and the discriminator:



$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x|y) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z, y|y)))]$$



# Let's Look at Some Code!

## Conditional GAN Example

Code on GitHub : [https://github.com/atapour/dl-pytorch/blob/main/Conditional\\_GAN\\_Example/Conditional\\_GAN\\_Example.ipynb](https://github.com/atapour/dl-pytorch/blob/main/Conditional_GAN_Example/Conditional_GAN_Example.ipynb)

Code on Colab: [https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Conditional\\_GAN\\_Example/Conditional\\_GAN\\_Example.ipynb](https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Conditional_GAN_Example/Conditional_GAN_Example.ipynb)

# Generative Adversarial Networks

Advanced  
Variants



- Auxiliary Classifier GAN (ACGAN)
- InfoGAN
- Least Square GAN (LSGAN)
- Adversarial Autoencoder
- Boundary Equilibrium GAN (BEGAN)
- Energy-Based GAN
- StyleGAN
- etc. etc. etc. ...

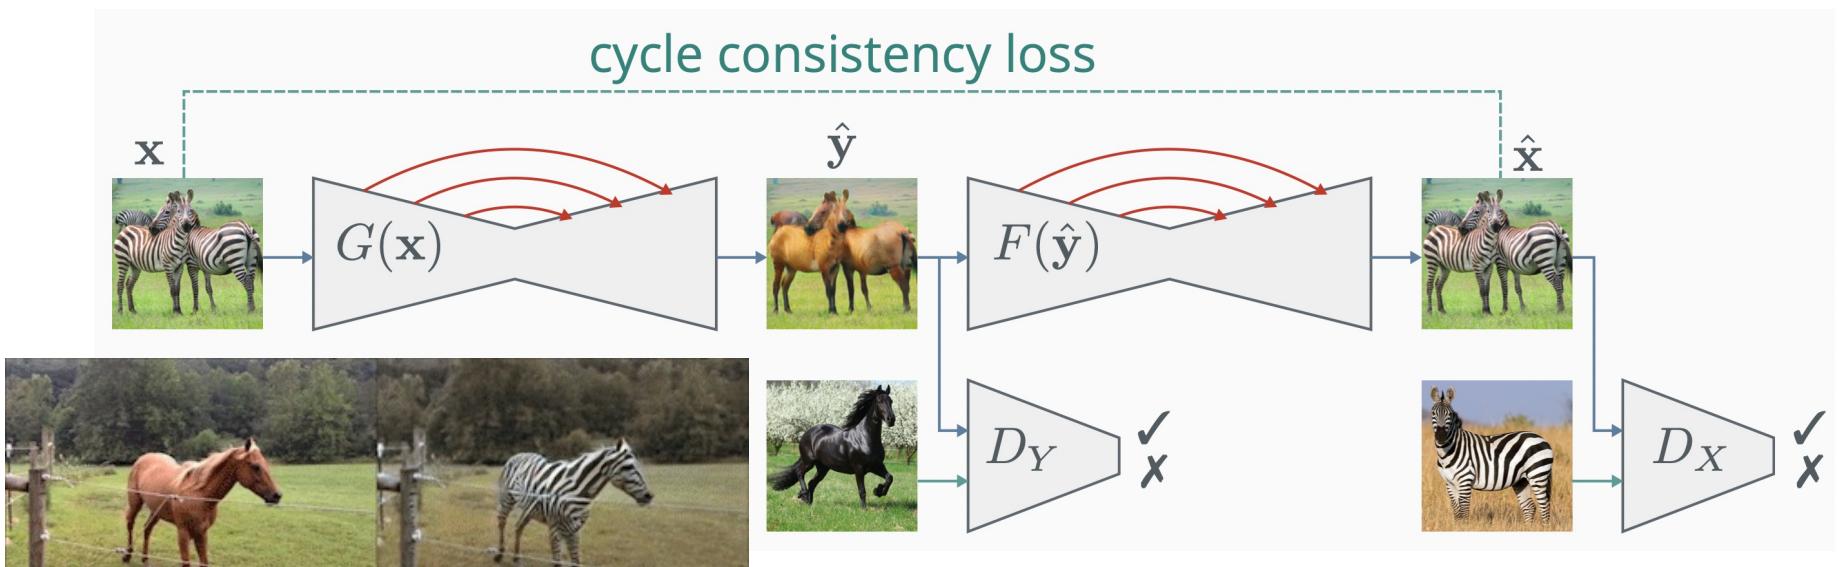
# Popular Applications

Unpaired Translation (CycleGAN)



**Definition:** CycleGAN [Zhu et al., 2017]

An adversarial architecture for unpaired image translation. It has twin generators with skip connections and two discriminators, which translate between the domains, alongside a cycle consistency loss (L1 distance) to ensure the mapping recovers the original image.





# Popular Applications

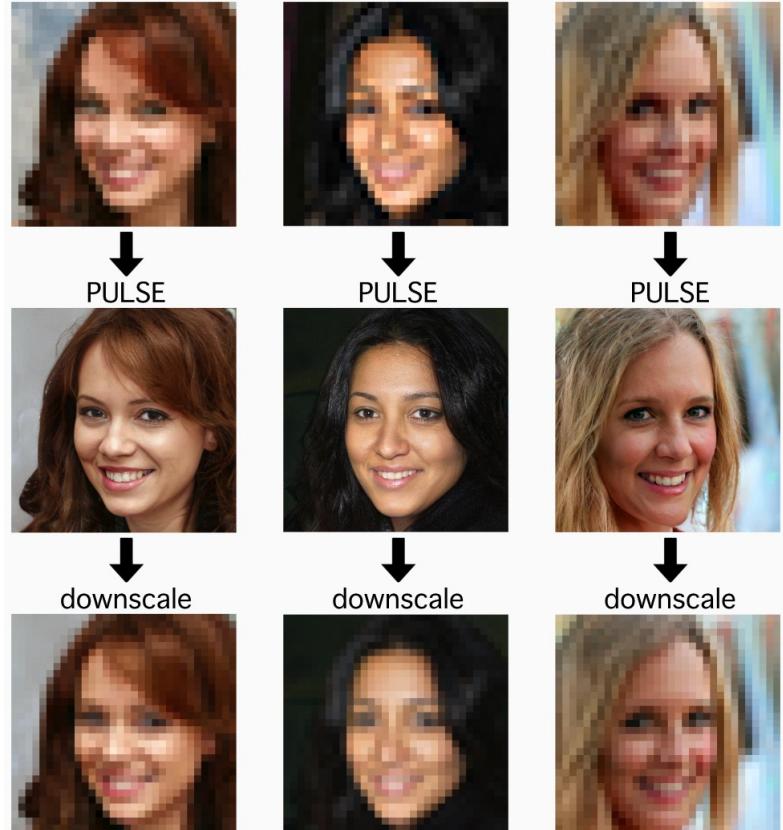
Super-resolution

## Definition: Super-resolution

The objective is to map a single low-resolution input to a distribution of high-resolution outputs. PULSE [Menon et al., 2020] projects points in the search of the latent space of StyleGAN (a large conditional GAN) onto a hypersphere, which ensures probable outputs in the high-dimensional latent space.



Pulse  
[\[demo\]](#)



# Popular Applications

## Anomaly Detection



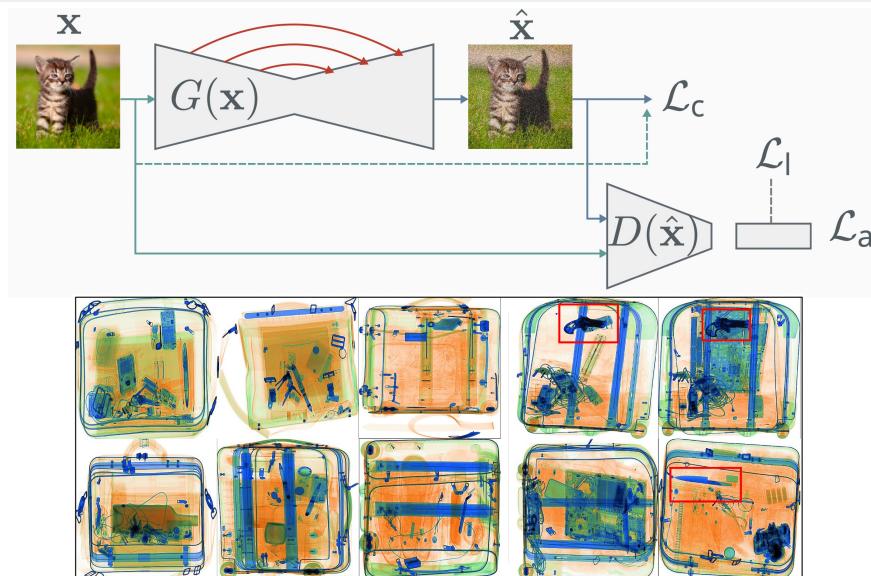
### Definition: Anomaly Detection

The objective is to find “anomalies” in the data. Unsupervised anomaly detectors [1,2] learn a normal distribution over (healthy) observations. Then, when they observe something not observed in training (unhealthy/dangerous), they fail to reconstruct - detecting it as an anomaly. Region-based anomaly detectors [3] learn a distribution over inpainted (erased) regions.

[1] Akçay, **Atapour-Abarghouei**, and Breckon. “GANomaly: Semi-Supervised Anomaly Detection via Adversarial Training”, ACCV, 2018.

[2] Akçay, **Atapour-Abarghouei**, and Breckon. “Skip-ganomaly: Skip connected and adversarially trained encoder-decoder anomaly detection”, IJCNN, 2019.

[3] Nguyen, Feldman, Bethapudi, Jennings, and Willcocks. “Unsupervised Region-based Anomaly Detection in Brain MRI with Adversarial Image Inpainting”, arXiv:2010.01942.





# What we learned today!

## 1 Generative models

- Unsupervised training
- PixelRNN and PixelCNN

## 2 Generative adversarial networks

- Definition
- Common Issues
- Lipschits Continuity
- Spectral Normalisation
- Conditional GANs
- Other Advanced Variants

## 3 Popular applications

- Unpaired Translation
- Super Resolution
- Anomaly Detection