

# Deep Learning

## Lecture 7: Sequential Models

---

Amir Atapour-Abarghouei

*amir.atapour-abarghouei@durham.ac.uk*

Durham University





# Lecture Overview

## 1 Recurrent neural networks

- Definition and implementation
- Backpropagation through time
- vanishing/exploding gradients

## 2 Long short-term memory

- Definition
- Properties
- Seq2Seq
- Attention

## 3 Transformers

- Self-Attention
- End-to-end object detection
- GPT
- DALL-E




# Recurrent Neural Networks

Definition

## Definition: recurrent neural networks


A function applied to nodes on a directed graph. *Unlike* one-way directed graphs (e.g. text, audio), sequential data is modelled using a cyclic connection that allows information to be stored [Rumelhart et al., 1986]. The same function  $f$  is applied to inputs at each time step, updating a hidden state vector  $h$  which acts as the network's memory:

$$h_{t+1} = f_\theta(h_t, x_t)$$






# Computational Graphs




One-to-One

*Feedforward network*




One-to-Many

*e.g., image captioning*



Many-to-One

*e.g., sentence classification*



Many-to-Many

*seq2seq*

# Recurrent Neural Networks




Implementation

## Example: RNN

A simple implementation is:

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

which is visually interpreted as a “cell”:



$$P(\mathbf{x}) = \prod_t P(x_t | x_{t-1}, x_{t-2}, \dots, x_1)$$

Output	e	l	l	o
Output Layer	1.0 2.2 -3.0 4.1	0.5 0.3 -1.0 1.2	0.1 0.5 1.9 -1.1	0.2 -1.5 -0.1 2.2
Hidden Layer	0.3 -0.1 0.9	1.0 0.3 0.1	0.1 -0.5 -0.3	-0.3 0.9 0.7
Input Layer	1 0 0 0	0 1 0 0	0 0 1 0	0 0 1 0
Input	h	e	l	l

# Recurrent Neural Networks

Backpropagation




## Definition:

Backpropagation applied to an **unrolled** RNN is called backprop through time (BPTT). Gradients accumulate in  $W$  additively:

$$\frac{\partial \mathcal{L}_T}{\partial W} = \sum_{t \leq T} \frac{\partial \mathcal{L}_T}{\partial h_t} \frac{\partial h_t}{\partial W}$$

Long sequences use truncated BPTT where sequences are split into batches but hidden connections remain.





# Let's Look at Some Code!




## Simple RNN Example for Text Classification

Code on GitHub : [https://github.com/atapour/dl-pytorch/blob/main/RNN\\_Sentiment\\_Analysis/RNN\\_Sentiment\\_Analysis.ipynb](https://github.com/atapour/dl-pytorch/blob/main/RNN_Sentiment_Analysis/RNN_Sentiment_Analysis.ipynb)

Code on Colab: [https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/RNN\\_Sentiment\\_Analysis/RNN\\_Sentiment\\_Analysis.ipynb](https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/RNN_Sentiment_Analysis/RNN_Sentiment_Analysis.ipynb)

# Recurrent Neural Networks

Exploding and Vanishing  
Gradients



Why do gradients vanish/explode?


The gradient of involves many factors of  $W$  (and tanh).

The product of  $T$  matrices converges to 0 (or grows to infinity) at an exponential rate in  $T$ .

$$\frac{\partial \mathcal{L}_T}{\partial W} = \sum_{t \leq T} \frac{\partial \mathcal{L}_T}{\partial h_t} \frac{\partial h_t}{\partial W} = \sum_{t \leq T} \frac{\partial \mathcal{L}_T}{\partial h_T} \frac{\partial h_T}{\partial h_t} \frac{\partial h_t}{\partial W}$$

[Demo](#)

Solution to exploding gradients: **clip gradients**



# Long Short-Term Memory

Preventing Vanishing Gradient



## Definition: long short term memory

LSTMs [Hochreiter et al., 1997] learn longer sequences than vanilla RNNs using better gradient flow. Backpropagation from  $c_t$  to  $c_{t-1}$  has no direct matrix multiplication by  $W$ .

**Gates** determine how much information passes through.

Always adding new information to the hidden state can be overwhelming.

Sometimes we want to forget things!

## Example: LSTM cell



Gates control how much information is passed through using a sigmoid function and a dot product.

# Let's Look at Some Code!



## Simple LSTM Example for Text Classification

Code on GitHub : [https://github.com/atapour/dl-pytorch/blob/main/LSTM\\_Sentiment\\_Analysis/LSTM\\_Sentiment\\_Analysis.ipynb](https://github.com/atapour/dl-pytorch/blob/main/LSTM_Sentiment_Analysis/LSTM_Sentiment_Analysis.ipynb)

Code on Colab: [https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/LSTM\\_Sentiment\\_Analysis/LSTM\\_Sentiment\\_Analysis.ipynb](https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/LSTM_Sentiment_Analysis/LSTM_Sentiment_Analysis.ipynb)

# Long Short-Term Memory



Properties

## LSTM Properties

### Main Strengths

- Allows for variable length sequences
- Efficient parameter usage
- Theoretically able to store arbitrarily old information

### Main Limitations

- Practically unable to store very long term dependencies
- Limited by fixed size of hidden state
- Slow training and synthesis

### Further Reading:

<https://karpathy.github.io/2015/05/21/rnn-effectiveness/>

# Seq2Seq Architecture

Natural Language Processing



- Seq2Seq (sequence to sequence) receives a sentence as the input and produces a sentence as the output.
- Seq2Seq will include two networks, one encoder and one decoder.
- An example of this would be “**Machine Translation**”:

Source Sentence (English)

*granny liked my dishes.*



Target Sentence (French)

*mamie a aimé mes mets.*






# Seq2Seq Architecture

Machine Translation

- Let's see how inference works. We will take about the training later.

The sentence representation is used as the initial hidden state for the decoder.






# Seq2Seq Architecture

Machine Translation

- Now, let's see how we would train this.

*Loss: cross-entropy*

- The system is trained "end to end" as one single entity.





# Seq2Seq Architecture


Issue

Solution: **Attention**

The entirety of the source sentence  
should be represented within the final  
hidden state of the encoder.


**Informational Bottleneck:**

too much pressure on a single vector






# Seq2Seq with Attention






# Seq2Seq with Attention






# Seq2Seq with Attention





# Seq2Seq with Attention





# Attention

## Definition and Calculation

1. Encoder hidden states:  $h_1, \dots, h_N \in \mathbb{R}^h$
2. At timestep  $t$ , decoder hidden state:  $s_t \in \mathbb{R}^h$
3. Attention scores at timestep  $t$ :  $e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$
4. Softmax to get attention distribution:  $\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$
5. Use  $\alpha^t$  to take a weighted sum of the encoder hidden states to get the attention output:  $a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$
6. Concatenate attention output  $a_t$  with the decoder state  $s_t$  and continue with the rest of model training:  $[a_t; s_t] \in \mathbb{R}^{2h}$

Encoder hidden states:  $h_1, \dots, h_N$

Decoder hidden state:  $s_t$

**Definition:** Given a set of vectors *values* and a vector *query*, attention is a weighted sum of the *values*, dependant on the *query*.



**Attention types:**  Dot product  Additive  Multiplicative



# Attention Is All You Need - Transformers

[Vaswani et al., 2017]

- **Self-Attention** provided a route to removing recurrence.
- **No RNNs.**




# Attention Is All You Need - Transformers



[Vaswani et al., 2017]

- **Self-Attention** provided a route to removing recurrence.
- **No RNNs.**




# Attention Is All You Need - Transformers



[Vaswani et al., 2017]

- **Self-Attention** provided a route to removing recurrence.
- **No RNNs.**






# Attention Is All You Need - Transformers

[Vaswani et al., 2017]

Parallelisable

Transformers offer improvements in performance and efficiency.



# Let's Look at Some Code!



## Simple Transformer Example for Text Classification

Code on GitHub : [https://github.com/atapour/dl-pytorch/blob/main/Transformer\\_Sentiment\\_Analysis/Transformer\\_Sentiment\\_Analysis.ipynb](https://github.com/atapour/dl-pytorch/blob/main/Transformer_Sentiment_Analysis/Transformer_Sentiment_Analysis.ipynb)

Code on Colab: [https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Transformer\\_Sentiment\\_Analysis/Transformer\\_Sentiment\\_Analysis.ipynb](https://colab.research.google.com/github/atapour/dl-pytorch/blob/main/Transformer_Sentiment_Analysis/Transformer_Sentiment_Analysis.ipynb)



# Transformers

GPT

## GPT Training and Evaluation

**GPT-3** [Brown et al., 2020]:

- 175B parameters (96 layers with 96 heads each with 12,228 neurons)
- Batch size 3.2M. Input length of 2048
- Terabytes of data from the internet

**GPT-4 and GPT-5:**

- **ChatGPT**: much larger

Evaluation Tasks

- Few shot translation
- Reading comprehension (Q&A)
- Closed book question & answering
- Natural language inference
- News article writing
- Dialogue generation

## Example: ChatGPT article



ChatGPT

Title: United Methodists Agree to Historic Split

Subtitle: Those who oppose gay marriage will form their own denomination

The United Methodist Church, one of the largest Protestant denominations in the United States, has reached a pivotal moment in its history. After years of internal conflict and debate over LGBTQ+ rights, particularly the issue of gay marriage and the ordination of LGBTQ+ clergy, the denomination has agreed to a historic split. This decision will see the formation of a new, conservative denomination by members who oppose gay marriage.

The schism is the culmination of a long-standing rift within the church, reflecting broader societal

## Alternative LLMs

- Google AI's Gemini
- Meta AI's LLaMA
- Anthropic's Claude
- Falcon
- ...



# Transformers

GPT-3: Codex

## Example: Codex

Codex is GPT trained on multiple datasets based on millions of GitHub repositories [one 2020 dataset consisted of 54 million public repos].

Available as an extension for Visual Studio Code and other IDEs.

VS Code Cursor + Sonnet 3.5

Output still contains bugs and errors but improving.

The screenshot shows the Visual Studio Code interface with the Python extension installed. The left sidebar has icons for files, search, and code navigation. The main editor window shows a file named 'main.py' with three lines of code: '1', '2', and '3'. The bottom right corner of the editor shows a small icon with 'co-pilot' next to it. The bottom navigation bar includes tabs for PROBLEMS, OUTPUT, TERMINAL, and DEBUG CONSOLE. The terminal tab is active, showing the command line: 'rahulbanerjee@Rahuls-MBP ~ %'. The status bar at the bottom provides information about the Python environment ('Python 3.9.4 64-bit (3.9)'), code analysis results ('0 ▲ 0'), and various development tools like Go Live, Chronicler, and Prettier.



# Transformers

DALL-E

## DALL-E Training


**DALL-E** [Ramesh et al., 2021]:

- 12-billion parameter version of GPT-3
- Generate images from detailed text descriptions
- **DALL·E 2** and now **DALL·E 3**

Capable of anthropomorphised versions of animals and objects, combining unrelated concepts in plausible ways, rendering text, and applying transformations to existing images

## Example: DALL-E Image

A robot teaching his grandmother to suck eggs.



Good-looking man with golden hair and big bushy beard.



# Transformers


## End-to-End Object Detection

### DEtection TRansformer

Fast object detection is crucial for many tasks including self-driving cars. Training end to end is difficult due to the discrete nature of objects.

DETR [Carion et al., 2020] enables global search and 'query' of the image for information. Attention matrices can also be used to make segmentation maps.

### Example: architecture and examples





# What we learned today!

## 1 Recurrent neural networks

- Definition and implementation
- Backpropagation through time
- vanishing/exploding gradients

## 2 Long short-term memory

- Definition
- Properties
- Seq2Seq
- Attention

## 3 Transformers

- Self-Attention
- End-to-end object detection
- GPT-3
- DALL-E