

Deep learning practical 1

Jupyter & Colab setup

Practical overview

Welcome to the first deep learning practical. This is only an initial introductory session and is only there to ensure a suitable deep learning training environment.

Environment setup

There are generally four ways to run the code and train the large deep learning models required for this module:

- Use Google Colab—easy to get started, but the usage gets throttled and the kernel is not persistent
- Use the NCC Jupyter server—recommended for most students. The I/O can be a little slow
- SSH to NCC and use SLURM
- Use the lab machines or your home machines—not recommended if the GPUs are inferior

Using Google Colab

- You need a Google account for this
- Simply open an existing .ipynb in Colab and run the cells—here are two to get started with:
 - Train a classifier: [🔗 download](#)
 - Train a convnet: [🔗 download](#)
- Check that your ‘Runtime > Change runtime type’ is a GPU

NCC usage via Jupyter (recommended)

- Connect to a machine from the university network, otherwise use the VPN to connect from outside
- Login to `nccl.clients.dur.ac.uk/COMP0000` with your CIS username and password (if you get an authentication failure, you might not have an account on NCC and need to email Rob Powell)
- Select appropriate preferences for the job, e.g. `ug-gpu-small` with 1 GPU
- Once it takes you to your home directory click “Assignments” tab and “Fetch” the environment creation folder
- Click the title of the folder “Create Jupyter Environment Script” and click the “CreateJupyterPythonEnvironment” script that is inside it
- In the second code cell you need to enter a Python version, and two names (these can be the same if you like). The `virtualEnvName` becomes the directory name of the `virtualenv` in your NCC home directory. The `jupKernelName` is what appears in the Jupyter menu later. If you’re only using PyTorch you shouldn’t need modules as CUDA is bundled
- Run the whole notebook (Kernel > Restart and Run All). Once it has finished (about 4-5 mins) you can close this tab and return to the previous tab that should still be on the Assignments page.
- Click “Files” to see home directory again and then refresh the page otherwise the new kernel won’t appear in the list!
- Click “New” and select your kernel name from the dropdown list
- Run your `!pip install` commands (only needs to be done the first time you ever use the kernel, unlike Colab, as our storage is persistent)
- Download, upload to NCC and run the `classifier.ipynb` [🔗](#) and `convnet.ipynb` examples [🔗](#)

NCC usage via SSH

Only do this if you know what you're doing, e.g. have used SLURM before:

- SSH to the head node `CIS-username@ncc1.clients.dur.ac.uk` (either via VPN or Mira)
 - Do not run code on this head node. Doing so will cause problems for all other users, can interrupt active postgraduate research and may get you banned
 - Before running any code, read the NCC documentation carefully: `ncc1.clients.dur.ac.uk` and do your compute with SLURM
 - I recommend using `tmux` for a persistent shell along with `zsh`
 - Please respect other users in the queue
-