

Reinforcement Learning

Lecture 2: Environments and bandits

Chris G. Willcocks

Durham University



This lecture shows how to set up a modern RL environment ready for training using **Gym** [↗](#), **Brax** [↗](#) and a multi-armed bandits implementation as in chapter 2 of [1].

1 Gym and Brax

- what are they?
- how do they map to RL concepts?
- gym.Env methods and attributes
- nested spaces example

2 Colab coding

- a random agent
- exploring different environments
- minimal REINFORCE example
- multi-armed bandits

3 Custom environments

- OpenAI gym
- Google brax



Gym and Brax what are they?

OpenAI gym & Google brax


What is it?

- Collection of environments
- Most environments are freely available
- Standardised API makes comparisons and benchmarking easier
- Test same RL algorithms on many different problems


For example:

- Atari 2600 games
- Classic control problems
- Physics simulations

Environments and performance in [2]



RL Agents



Algorithm: minimal example

```
import gym
env = gym.make('CartPole-v0')

agent = Agent()

# init the environment to base state
# and get initial observation state
s = env.reset()

# loop until environment says it's done
done = False
while not done:
    a = agent.sample_action(s)
    s, r, done, info = env.step(a)
env.close()
```



Core methods

In gym/core.py:

- **s = env.reset()** - reset the environment and return the initial observation state
- **s,r,done,info = env.step(action)** - takes a single step and returns a tuple with the next state, the reward, whether the environment is now done, and some meta diagnostic info
- **env.close()** - cleanup function
- **env.seed(int)** - seed all the environment randomness
- **env.render()** - (optional) render the env state: print text, show an image...

Env attributes

- **env.action_space** - gym.Space object that describes the shape and type of actions
- **env.observation_space** - gym.Space object that describes the shape and type of observations
- **env.reward_range** - tuple of per-step reward range, defaults to $(-\infty, \infty)$

gym.Space in /gym/spaces/space.py:

- **Box, Discrete, Binary, ...**
- **Tuple, Dict** - can contain nested observation spaces



Example: nested observation space – gym/spaces/dict.py

```
env.nested_observation_space = spaces.Dict({
    'sensors': spaces.Dict({
        'position': spaces.Box(low=-100, high=100, shape=(3,)),
        'velocity': spaces.Box(low=-1, high=1, shape=(3,)),
        'front_cam': spaces.Tuple((
            spaces.Box(low=0, high=1, shape=(10, 10, 3)),
            spaces.Box(low=0, high=1, shape=(10, 10, 3))
        )),
        'rear_cam': spaces.Box(low=0, high=1, shape=(10, 10, 3)),
    }),
    'ext_controller': spaces.MultiDiscrete((5, 2, 2))
})
```



Click this link:

Link to Colab code 

In this demo we will cover:

- A random agent
- Exploring different environments
- Plotting and measuring baselines
- Minimal REINFORCE example
- Multi-armed Bandits – chapter two of [1]



Example: gym custom environment

```
class CustomEnv(gym.Env):
    def __init__(self):
        ...

    def _seed(self, seed=None):
        ...

    def step(self, action):
        ...
        return state, reward, done, info

    def reset(self):
        return state

    def render(self):
        pass
```





Example: brax custom environment [3]

```
from brax.envs import env
class Fast(env.Env):
    def __init__(self, **kwargs):
        super().__init__(config='dt: .02', **kwargs)

    def reset(self, rng: jnp.ndarray) -> env.State:
        zero = jnp.zeros(1)
        qp = brax.QP(pos=zero, vel=zero, rot=zero, ang=zero)
        obs, reward, done = jnp.zeros(2)
        return env.State(qp, obs, reward, done)

    def step(self, state: env.State, action: jnp.ndarray) -> env.State:
        vel = state.qp.vel + (action > 0) * self.sys.config.dt
        pos = state.qp.pos + vel * self.sys.config.dt
        qp = state.qp.replace(pos=pos, vel=vel)
        obs = jnp.array([pos[0], vel[0]])
        reward = pos[0]
        return state.replace(qp=qp, obs=obs, reward=reward)
```



- [1] Richard S Sutton and Andrew G Barto.
Reinforcement learning: An introduction (second edition). Available online  MIT press, 2018.
- [2] Volodymyr Mnih et al. "Human-level control through deep reinforcement learning". In: nature 518.7540 (2015), pp. 529–533.
- [3] C Daniel Freeman et al. "Brax–A Differentiable Physics Engine for Large Scale Rigid Body Simulation". In: arXiv preprint arXiv:2106.13281 (2021).