

# Reinforcement Learning

## Lecture 4: Dynamic programming

---

Robert Lieck

Durham University

Lecture covers Chapter 4 in Sutton & Barto [1] and adaptations from David Silver [2]

## 1 Introduction

---

- definition
- examples
- planning in an MDP

## 2 Policy evaluation

---

- definition
- synchronous algorithm

## 3 Policy iteration

---

- policy improvement
- definition
- modified policy iteration

## 4 Value iteration

---

- definition
- summary and extensions

## Definition: Dynamic programming

Dynamic programming is an optimisation method for sequential problems. DP algorithms are able to solve complex 'planning' problems.

Given a complete MDP, dynamic programming can find an optimal policy. This is achieved with two principles:

1. Breaking down the problem into subproblems
2. Caching and reusing optimal solutions to subproblems to find the overall optimal solution

## Planning: what's the optimal policy?






## Famous examples

- Dijkstra's algorithm
- Backpropagation
- Doing basic math

...so it's really just recursion and common sense!





## Dynamic programming for planning MDPs

In reinforcement learning, we want to use dynamic programming to solve MDPs. So given an MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$ :

First, we want to find the value function  $v_\pi$  for that policy:

- This is done by **policy evaluation** (the prediction problem)

Then, when we're able to evaluate the policy, we want find the best policy  $v_*$  (the control problem). This is done with two strategies:

1. **Policy iteration**
2. **Value iteration**

**Follow along in notebook:**

## Definition: Policy evaluation

We want to evaluate a given policy  $\pi$ .  
We'll achieve this with the Bellman **expectation** equation,  $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$

	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

## Example: frozen lake environment





## Algorithm: policy evaluation

```
def policy_evaluation(env, policy,  $\gamma$ , theta):  
     $\rightarrow V = \text{np.zeros}(\text{env.num\_states})$   
    while True:  
        delta = 0  
        for s in range(env.num_states):  
            Vs = 0  
            for a, a_prob in enumerate(policy[s]):  
                for prob, s', reward, done in env.P[s][a]:  
                    Vs += a_prob * prob * (reward +  $\gamma$  * V[s'])  
            delta = max(delta, abs(V[s]-Vs))  
            V[s] = Vs  
        if delta < theta:  
            break  
    return V
```

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	



## Recap: Bellman expectation equation

$$\begin{aligned}v_{\pi}(s) &= \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a) \\&= \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s') \right)\end{aligned}$$

## Algorithm: policy evaluation

```
(iteration=1,  $\gamma=1$ )
for s in range(env.num_states):
    Vs = 0
    for a, a_prob in enumerate(policy[s]):
        for prob, s', reward, done in env.P[s][a]:
            Vs += a_prob * prob * (reward +  $\gamma$  * V[s'])
→ V[s] = Vs
```

iteration 1,  $\pi = \begin{smallmatrix} \nearrow & \nwarrow \\ \swarrow & \searrow \end{smallmatrix}$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.25	





## Recap: Bellman expectation equation

$$\begin{aligned}v_{\pi}(s) &= \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a) \\&= \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s') \right)\end{aligned}$$

## Algorithm: policy evaluation

```
(iteration=2,  $\gamma=1$ )  
for s in range(env.num_states):  
    Vs = 0  
    for a, a_prob in enumerate(policy[s]):  
        for prob, s', reward, done in env.P[s][a]:  
            Vs += a_prob * prob * (reward +  $\gamma$  * V[s'])  
→ V[s] = Vs
```

iteration 2,  $\pi = \begin{smallmatrix} \nearrow & \nwarrow \\ \swarrow & \searrow \end{smallmatrix}$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.06	0.0
0.0	0.06	0.34	



iteration 3,  $\pi = \begin{smallmatrix} \nearrow & \nwarrow \\ \swarrow & \searrow \end{smallmatrix}$

		.016	
	.031	.098	
	.109	.388	

iteration 4,  $\pi = \begin{smallmatrix} \nearrow & \nwarrow \\ \swarrow & \searrow \end{smallmatrix}$

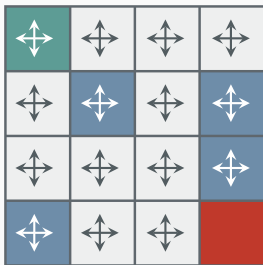
		.004	.001
		.025	
.008	.054	.117	
	.138	.411	

...

iteration  $\infty$ ,  $\pi = \begin{smallmatrix} \nearrow & \nwarrow \\ \swarrow & \searrow \end{smallmatrix}$

.014	.012	.021	.010
.016		.041	
.035	.088	.142	
	.176	.439	

random policy

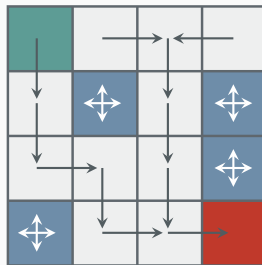


iteration  $\infty$ ,  $\pi =$  

.014	.012	.021	.010
.016		.041	
.035	.088	.142	
	.176	.439	

$\max_a$

improved policy



## Definition: Policy iteration

Given a policy  $\pi$  (e.g. starting with a random policy), **iteratively** evaluate:

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots \mid S_t = s]$$
$$\pi' = \text{greedy}(v_{\pi})$$

This always converges to the optimal policy  $\pi^*$ . That is, if the improvements stop:

$$q_{\pi}(s, \pi'(s)) = \max_{a \in \mathcal{A}} q_{\pi}(s, a) = q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

then the Bellman equation has been satisfied  $v_{\pi}(s) = \max_{a \in \mathcal{A}} q_{\pi}(s, a)$  therefore  $v_{\pi} = v_*(s)$  for all  $s \in \mathcal{S}$

## Example: learning a better policy





## Algorithm: modified policy iteration

What if we don't do iterative policy evaluation to  $\infty$ ?  
What if we just do a crude, e.g.  $k = 3$  small amount of iteration?

Does it still converge?

- Yes! It still converges to the optimal policy
- except in the case  $k = 1$  which is equivalent to value iteration

iteration 3,  $\pi =$  

		.016	
	.031	.098	
	.109	.388	



## Bellman optimality equation

If we recap the definition of the optimal value function according to the Bellman optimality equation:

$$\begin{aligned} v_*(s) &= \max_a q_*(s, a) \\ &= \max_a \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s') \end{aligned}$$

We can also iteratively apply the update with the one-step look-ahead to learn  $v_*(s)$

## Algorithm: value iteration

```
def value_iteration(env,  $\gamma$ , theta):  
    V = np.zeros(env.nS)  
    while True:  
        delta = 0  
        for s in range(env.nS):  
            v_s = V[s]  
            q_s = np.zeros(env.nA)  
            for a in range(env.nA):  
                for prob, s', reward, done in env.P[s][a]:  
                    q_s[a] += prob * (reward +  $\gamma$  * V[s'])  
            V[s] = max(q_s)  
            delta = max(delta, abs(V[s] - v_s))  
        if delta < theta: break  
    policy = greedily_from(env, V, gamma)  
    return policy, V
```




## Summary


In summary, dynamic programming:

- solves the planning problem, but not the full reinforcement learning problem
- requires a complete model of the environment
- policy evaluation solves the prediction problem
- there's a spectrum between policy iteration and value iteration
- these solve the control problem

Extensions:

- Asynchronous DP (read section 4.5 of Sutton & Barto [1])
- Play with the [interactive demo by Andrej Karpathy](#) 



- [1] Richard S Sutton and Andrew G Barto.  
Reinforcement learning: An introduction (second edition). Available online . MIT press, 2018.
- [2] David Silver. Reinforcement Learning lectures.  
<https://www.davidsilver.uk/teaching/>. 2015.