

# Reinforcement Learning

## Lecture 10: Extended methods

---

Robert Lieck

Durham University



## 1 The future of RL

---

- The trilemma in RL/ML/AI
- The bitter lesson
- Reward is (not) enough?

## 2 Long-term dependencies

---

- DQN characteristics
- Recurrent and distributed RL
- R2D2 performance
- Representation learning

## 3 Intrinsic reward

---


- Exploration vs exploitation

## 4 Prior knowledge

---

- AlphaStar and StarCraft
- meta-knowledge and training heuristics

# The trilemma in RL/ML/AI



## Extreme strategies


- if you understand your problem  
→ **hard-code** the solution
- if you have enough computation power  
→ **brute-force** search the solution space
- if you have enough data (and computation power)  
→ learn a general function approximator to **interpolate** between the data points

## Long-term strategy/development



1. use prior knowledge to simplify problem if needed
2. replace *hard* and *specific* assumptions by *soft* and *general* heuristics
3. progressively move towards general-purpose solutions

## In reality...

...we need to combine these strategies. We cannot (efficiently) learn anything without making assumptions (prior knowledge) [1].




## Rich Sutton, 2019

 Rich Sutton's bitter lesson is that, despite it being tempting to incorporate domain knowledge, in the long run, general purpose agents win. [Link to article](#) 

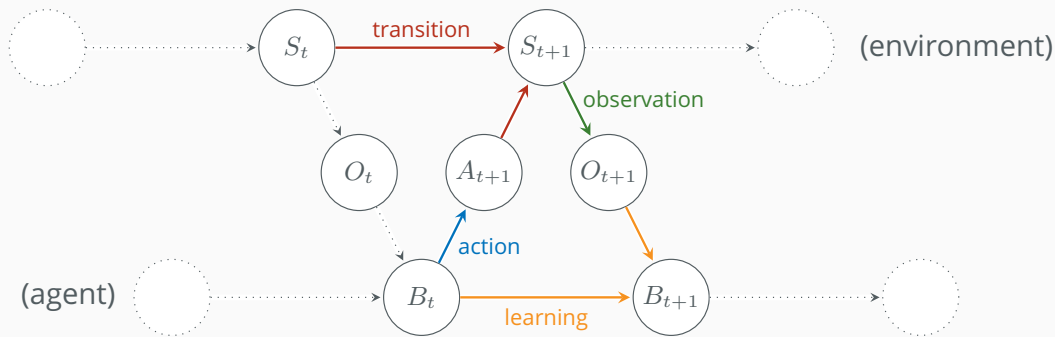
- AI researchers have often tried to build knowledge into their agents
- this always helps in the short term and is personally satisfying to the researcher
- but in the long run it plateaus and even inhibits further progress
- breakthrough progress eventually arrives by an opposing approach based on scaling computation by search and learning

# Reward is (not) enough?

Reward is enough [2] (Silver, Singh, Precup & Sutton). Others argue for intrinsic rewards in practice.




- For **defining** a problem?
- For **solving** a problem?



- $S$  is not known  $\rightarrow$  process is not Markov (w.r.t. the observations  $O$ )
- predictions depend on the entire history of actions and observations

## Characteristics: DQN

DQNs optimise a function (neural network) to predict the  $Q$ -value (the expected reward) for a given state and action.



DQN doesn't work very well for long-term credit assignments:



## Recap: function approximation

There are too many states/actions to fit into memory, so we estimate the value function:

$$\hat{v}(S, \mathbf{w}) \approx v_{\pi}(S),$$

or for control we'd do:

$$\hat{q}(S, A, \mathbf{w}) \approx q_{\pi}(S, A),$$

This usually requires several extra tricks:


- Double DQN
- Prioritised experience replay buffer
- Noisy linear layers



## Definition: R2D2

Recurrent Replay Distributed DQN (R2D2) [3] uses RNNs, training on a sequence of  $m = 80$  observations  $o_t$  and hidden states  $h_t$ :


Computation of  $\Delta Q$




Therefore it can backpropagate through the history, updating where earlier actions led to long-term future reward.

## Definition: distributed RL

In distributed RL [4], a central learner (with some parameters  $\theta$ ) receives experience from multiple parallel workers  $w_1, w_2, \dots, w_n$  which run episodes independently:




These graphs shows R2D2 performance for different ablations:




[Watch R2D2 play Gravitar](#)  [Watch R2D2 play other Atari](#) 


Dreamer [5] and DreamerV2 [6] use a recurrent neural network to ‘imagine’ and plan ahead, all in the latent (feature representation) space:



(a) Learn dynamics from experience



(b) Learn behavior in imagination



(c) Act in the environment



# Representation learning Dreamer algorithm

Initialize dataset  $\mathcal{D}$  with  $S$  random seed episodes. Initialize neural network parameters  $\theta, \phi, \psi$  randomly.

**while not converged do**

**for update step  $c = 1..C$  do**

// Dynamics learning

Draw  $B$  data sequences  $\{(a_t, o_t, r_t)\}_{t=k}^{k+L} \sim \mathcal{D}$ .

Compute model states  $s_t \sim p_\theta(s_t | s_{t-1}, a_{t-1}, o_t)$ .

Update  $\theta$  using representation learning.

// Behavior learning

Imagine trajectories  $\{(s_\tau, a_\tau)\}_{\tau=t}^{t+H}$  from each  $s_t$ .

Predict rewards  $E(q_\theta(r_\tau | s_\tau))$  and values  $v_\psi(s_\tau)$ .

Compute value estimates  $V_\lambda(s_\tau)$  via [Equation 6](#).

Update  $\phi \leftarrow \phi + \alpha \nabla_\phi \sum_{\tau=t}^{t+H} V_\lambda(s_\tau)$ .

Update  $\psi \leftarrow \psi - \alpha \nabla_\psi \sum_{\tau=t}^{t+H} \frac{1}{2} \|v_\psi(s_\tau) - V_\lambda(s_\tau)\|^2$ .

// Environment interaction

$o_1 \leftarrow \text{env.reset}()$

**for time step  $t = 1..T$  do**

Compute  $s_t \sim p_\theta(s_t | s_{t-1}, a_{t-1}, o_t)$  from history.

Compute  $a_t \sim q_\phi(a_t | s_t)$  with the action model.

Add exploration noise to action.

$r_t, o_{t+1} \leftarrow \text{env.step}(a_t)$ .

Add experience to dataset  $\mathcal{D} \leftarrow \mathcal{D} \cup \{(o_t, a_t, r_t)_{t=1}^T\}$ .

## Model components

Representation  $p_\theta(s_t | s_{t-1}, a_{t-1}, o_t)$

Transition  $q_\theta(s_t | s_{t-1}, a_{t-1})$

Reward  $q_\theta(r_t | s_t)$

Action  $q_\phi(a_t | s_t)$

Value  $v_\psi(s_t)$

## Hyper parameters

Seed episodes  $S$

Collect interval  $C$

Batch size  $B$

Sequence length  $L$

Imagination horizon  $H$

Learning rate  $\alpha$

## Exploration vs exploitation

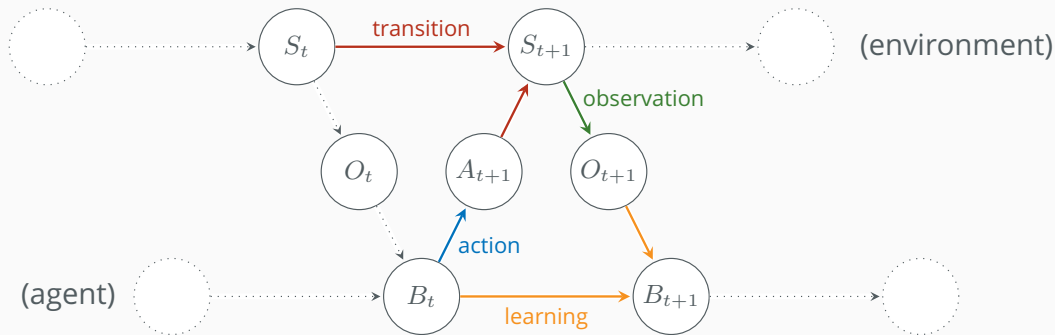
We can take random actions for exploration, but there are other strategies

- random exploration (as before)
  - $\epsilon$ -greedy
  - softmax
- optimism in the face of uncertainty
  - estimate uncertainty of the value
  - prefer exploring states/actions with higher uncertainty
- information/belief state space
  - the agent information is part of the state description
  - quantifies state information value

## Exploration in Gravitar and AoE

Randomly choosing isn't always good:






- $S$  is not known  $\rightarrow$  plan in belief space  $B$  instead
- uncertainty can be quantified and taken into account
- agent can actively plan for gathering useful information (*active learning*)
- hard to solve exactly  $\rightarrow$  use approximations and Monte-Carlo methods

AlphaStar [7] uses many components, supervised learning, and league-play.





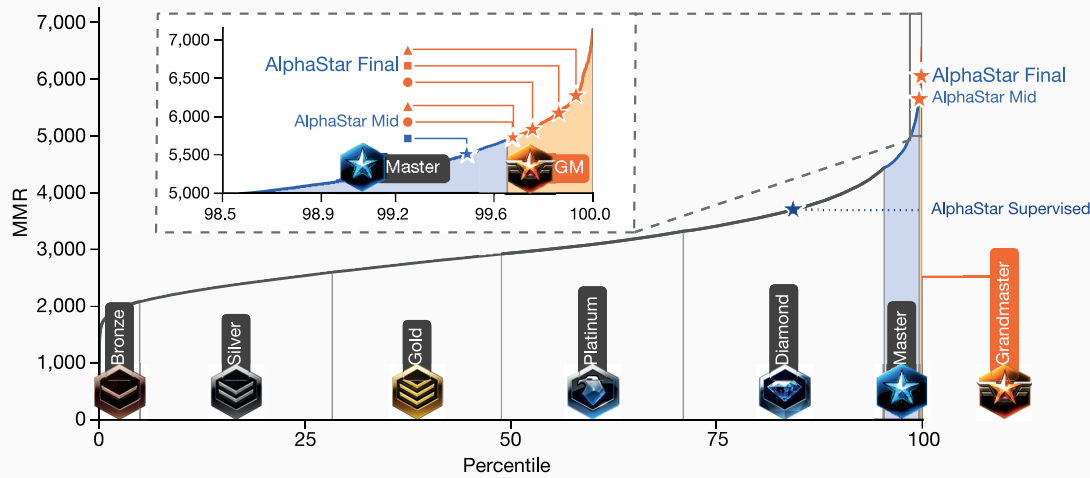
*“We found our use of human data to be critical in achieving good performance with reinforcement learning” [7]*



- start with training agent from human data (supervised learning)
- intrinsic rewards for following human statistics  $z$
- self-play and league-play
- policy gradient similar to actor-critic
- temporal difference learning (TD( $\lambda$ ))
- clipped importance sampling (V-trace)
- *new* self-imitation algorithm (UPGO)
- use opponent's observations during training
- league training
  1. **main agents:** prioritized fictitious self-play (PFSP) against mixture of agents
  2. **main exploiters:** play only against main agents to find weaknesses
  3. **league exploiters:** also PFSP, but are not targeted by main exploiters



# Prior knowledge AlphaStar results





## Summary

In summary:

- learn the foundations and concepts of the field, so you can speak the lingo...
- ...but you may want to approach overly complex papers more like an engineer
  - run the code and dismantle it back down to the concepts that make it work
- sample efficiency is an issue, which can be traded for with model-based imagination
- general purpose agents are the future



- [1] D. H. Wolpert and W. G. Macready. "No Free Lunch Theorems for Optimization". In: IEEE Transactions on Evolutionary Computation (1997).
- [2] David Silver et al. "Reward is enough". In: Artificial Intelligence (2021), p. 103535.
- [3] Steven Kapturowski et al. "Recurrent experience replay in distributed reinforcement learning". In: International Conference on Learning Representations. 2018.
- [4] Dan Horgan et al. "Distributed Prioritized Experience Replay". In: International Conference on Learning Representations. 2018.
- [5] Danijar Hafner et al. "Dream to Control: Learning Behaviors by Latent Imagination". In: International Conference on Learning Representations. 2020.
- [6] Danijar Hafner et al. "Mastering Atari with Discrete World Models". In: International Conference on Learning Representations. 2021.
- [7] Oriol Vinyals et al. "Grandmaster level in StarCraft II using multi-agent reinforcement learning". In: Nature 575.7782 (2019), pp. 350–354.