

COMP3667: Reinforcement learning practical 6

Temporal-Difference Learning

robert.lieck@durham.ac.uk

1 Overview

Welcome to the sixth reinforcement learning practical. In this practical, we will be experimenting with different TD methods to better understand their characteristics, advantages, and drawbacks. In particular, we will

- evaluate on-policy SARSA(0) and off-policy Q -Learning
- experiment with n -step TD learning
- compare the performance of these methods in different environments
- understand when/how/why they may fail to work and how to fix them.

2 Setup

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import itertools
4 from IPython import display
5 import rldurham as rld
6 from rldurham import plot_frozenlake as plot
```

As before, we will use different versions of the frozen lake gym environment:

```
1 name = 'FrozenLake-v1'
2 env = rld.make(name, is_slippery=False) # 4x4
3 env = rld.make(name, map_name="8x8", is_slippery=False) # 8x8
4 # env = rld.make(name, desc=["SFHH",
5 #                             "HFFH",
6 #                             "HHFF",
7 #                             "HHHG",], is_slippery=False) # custom
8 rld.seed_everything(42, env)
9 # LEFT, DOWN, RIGHT, UP = 0, 1, 2, 3
```

You can use these two helper classes to define hard-coded policies or policies using Q -values

```
1 class QPolicy:
2     def __init__(self, Q, epsilon):
3         self.Q = Q
4         self.epsilon = epsilon
5
6     def sample(self, state):
7         return np.random.choice(np.arange(self.Q.shape[1]), p=self.Q[state])
8         if np.random.rand() > self.epsilon:
9             best_actions = np.argwhere(self.Q[state]==np.max(self.Q[state])).flatten()
10            return np.random.choice(best_actions)
11        else:
12            return env.action_space.sample()
13
14    def __getitem__(self, state):
```

```

15     Qs = self.Q[state]
16     p = np.zeros_like(Qs)
17     max_actions = np.argmax(Qs == Qs.max())
18     p[max_actions] = 1 / len(max_actions)
19     return (1 - self.epsilon) * p + self.epsilon / len(p)

```

We can keep some plotting data in these variables (re-evaluate the cell to clear data)

```

1 reward_list = [[]]
2 auc = [0]
3 test_reward_list = [[]]
4 test_auc = [0]
5 plot_data = [[]]
6 plot_labels = []
7 experiment_id = 0

```

and use these functions to update and plot the learning progress

```

1 # (using global variables in functions)
2 def update_plot(mod):
3     reward_list[experiment_id].append(reward_sum)
4     auc[experiment_id] += reward_sum
5     test_reward_list[experiment_id].append(test_reward_sum)
6     test_auc[experiment_id] += test_reward_sum
7     if episode % mod == 0:
8         plot_data[experiment_id].append([episode,
9                                         np.array(reward_list[experiment_id]).mean(),
10                                         np.array(test_reward_list[experiment_id]).mean()])
11     reward_list[experiment_id] = []
12     test_reward_list[experiment_id] = []
13     for i in range(len(plot_data)):
14         lines = plt.plot([x[0] for x in plot_data[i]],
15                         [x[1] for x in plot_data[i]], '-',
16                         label=f"{plot_labels[i]}, AUC: {auc[i]}|{test_auc[i]}")
17         color = lines[0].get_color()
18         plt.plot([x[0] for x in plot_data[i]],
19                 [x[2] for x in plot_data[i]], '--', color=color)
20     plt.xlabel('Episode number')
21     plt.ylabel('Episode reward')
22     plt.legend()
23     display.clear_output(wait=True)
24     plt.show()
25
26 def next_experiment():
27     reward_list.append([])
28     auc.append(0)
29     test_reward_list.append([])
30     test_auc.append(0)
31     plot_data.append([])
32     return experiment_id + 1

```

3 On-policy and off-policy learning with TD(0)

Recap TD Learning

Remember our 0-step temporal difference (TD) targets from the lecture, which can be computed for any (also partial) episodes

$$\begin{aligned}
 V_{\pi}(s_t) &= \sum_{a_t \in \mathcal{A}} \pi(a_t | s_t) \sum_{s_{t+1} \in \mathcal{S}} p(s_{t+1} | s_t, a_t) \left[\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma V_{\pi}(s_{t+1}) \right] \\
 &= \mathbb{E}_{a_t \sim \pi(a_t | s_t)} \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} \left[\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma V_{\pi}(s_{t+1}) \right] \\
 &\approx \underbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma V_{\pi}(s_{t+1})}_{\text{TD(0) target}} \\
 Q_{\pi^*}(s_t, a_t) &= \sum_{s_{t+1} \in \mathcal{S}} p(s_{t+1} | s_t, a_t) \left[\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma \sum_{a_{t+1} \in \mathcal{A}} \pi^*(a_{t+1} | s_{t+1}) Q_{\pi^*}(s_{t+1}, a_{t+1}) \right] \\
 &= \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} \left[\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma \mathbb{E}_{a_{t+1} \sim \pi^*(a_{t+1} | s_{t+1})} Q_{\pi^*}(s_{t+1}, a_{t+1}) \right] \\
 &\approx \underbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma Q_{\pi^*}(s_{t+1}, a_{t+1})}_{\text{TD(0) target}} .
 \end{aligned}$$

At a particular time t we are in state s_t and take action $a_t \sim \pi(a_t | s_t)$ sampled from the policy π . We then end up in state $s_{t+1} \sim p(s_{t+1} | s_t, a_t)$ based on the environment's transition function. If we are interested in learning state-action values (i.e. solving the control problem), we additionally need to consider the following action $a_{t+1} \sim \pi^*(a_{t+1} | s_{t+1})$ based on the policy π^* that we want to learn or evaluate. Note that the sampling policy π and the policy π^* we want to evaluate are the same in *on-policy* methods but may be different in *off-policy* methods.

The TD(0) targets are noisy “snapshots” of how the values should look like based on the current transition at time t . The difference between the TD target and our current value estimate gives us a noisy TD error signal that tells us “how far off” our estimates are. This can be used to update our value estimates with a learning rate α (similar to SGD) to improve them

$$\begin{aligned}
 V_{\pi}(s_t) &\leftarrow V_{\pi}(s_t) + \alpha \left[\underbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma V_{\pi}(s_{t+1})}_{\text{TD(0) target}} - V_{\pi}(s_t) \right] \\
 Q_{\pi^*}(s_t, a_t) &\leftarrow Q_{\pi^*}(s_t, a_t) + \alpha \left[\underbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma Q_{\pi^*}(s_{t+1}, a_{t+1})}_{\text{TD(0) target}} - Q_{\pi^*}(s_t, a_t) \right] .
 \end{aligned}$$

On-policy SARSA(0)

We can evaluate and improve our policy “on the go”. This is called *on-policy* learning and it means that the policy π we use for sampling is the same as the policy π^* we are evaluating and learning.

Exercise

- Implement SARSA(0) by filling in the TD(0) targets and updates in the skeleton below. (The rest of the skeleton is for collecting episodes, evaluating the learned policy and plotting everything. The dashed line is the reward for the learned policy, the solid line is for the sampling policy.)
- Run a couple of evaluations with different values for the learning rate **alpha** and the exploration **epsilon** on the 4x4 environment. Use both “noisy” and “neutral” initialisations for Q (commenting in/out the respective lines in the code). What effects do you observe? With what parameters does the agent learn best?

You can use the following code as a basis and fill in the blanks:

```

1  # parameters
2  num_episodes = 3000
3  alpha = 0.1
4  gamma = 0.9
5  epsilon = 0.5
6  on_policy = True # SARSA or Q-Learning

```



```

67     # testing the learned policy
68     state, _ = env.reset()
69     test_reward_sum = 0
70     while True:
71         action = learned_policy.sample(state)
72         next_state, reward, term, trun, _ = env.step(action)
73         done = term or trun
74         test_reward_sum += reward
75         state = next_state
76         if done:
77             break
78
79     update_plot(int(np.ceil(num_episodes / 20)))
80
81 experiment_id = next_experiment()
82 print("Sampling policy and values")
83 plot(env, v=V, policy=sample_policy, draw_vals=True)
84 print("Learned policy and optimal/max values")
85 plot(env, v=Q.max(axis=1), policy=learned_policy, draw_vals=True)

```

Off-policy Q-Learning

Sometimes, we would like to generate samples with one policy π (typically an exploratory policy) but then use the samples to learn another policy π^* (typically a near-optimal policy). This is possible with *off-policy* methods, such as Q-Learning.

3.0.1 Exercise

- Test Q-Learning by using a different policy for the `learned_policy` (with lower `epsilon`) then for the `sample_policy`.
 - *Note:* Strictly speaking, the `sample_policy` does not change in Q-Learning. You can achieve this by using `epsilon=1` in the `sample_policy` to select actions randomly. For other values of `epsilon`, the `sample_policy` “peeks” into the values of the learned policy, which lets it profit from that learning (but is not a clean implementation).
 - *Bonus exercise:* Learn the `sample_policy` using SARSA(0) while learning the `learned_policy` with Q-Learning (this requires maintaining two copies of values estimates, performing separate updates for both etc).
- Run the evaluation several times on the 4x4 and 8x8 environment. Try to get the fastest and most reliable learning by tweaking `epsilon`. How does the performance (episode reward) of the sampling policy (solid line) compare to that of the learned policy (dashed lines)? How is this different from SARSA(0)?

Expected TD(0) targets

The normal TD targets to estimate state-action values are computed by *sampling* from the policy π^* that is to be learned. This sampling step increases the noise in the TD error signal (in addition to the noise we already have due to sampling episodes). Instead, one can use the *expected* TD targets

$$\underbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma \sum_{a_{t+1} \in \mathcal{A}} \pi^*(a_{t+1} | s_{t+1}) Q_{\pi^*}(s_{t+1}, a_{t+1})}_{\text{expected TD(0) target}} .$$

Exercise

- In the `q_target` replace the sampled value with an expectation over possible actions the `learned_policy` could take. *Note:* `learned_policy[next_state]` gives you the action probabilities for `next_state`.
- Compare the performance when using the expected TD(0) target to using the sampled one.

4 TD(n)

In TD(0) we do not look ahead and receiving a reward does only affect the value estimate of the current state and action. All the future expected rewards are approximated by using the current value estimates. However, we can improve on that by instead look n steps ahead (in hindsight) and taking the n next steps into account for computing TD targets. These TD(n) targets are

$$\begin{aligned} V_{\pi}(s_t) &\approx \overbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma \mathcal{R}_{s_{t+1} s_{t+2}}^{a_{t+1}} + \dots + \gamma^n \mathcal{R}_{s_{t+n} s_{t+n+1}}^{a_{t+n}}}^{\text{n-step return}} + \gamma^{n+1} V_{\pi}(s_{t+n+1}) \\ &\quad \text{TD}(n) \text{ target} \\ Q_{\pi^*}(s_t, a_t) &\approx \underbrace{\mathcal{R}_{s_t s_{t+1}}^{a_t} + \gamma \mathcal{R}_{s_{t+1} s_{t+2}}^{a_{t+1}} + \dots + \gamma^n \mathcal{R}_{s_{t+n} s_{t+n+1}}^{a_{t+n}}}_{\text{TD}(n) \text{ target}} + \gamma^{n+1} Q_{\pi^*}(s_{t+n+1}, a_{t+n+1}) . \end{aligned}$$

To compute them in hindsight, we have to store a trace of the last $n + 1$ transitions (including the current transition that is also used in $\text{TD}(0)$). The modified skeleton below is providing a trace of length $n + 1$ containing the last n steps and the current transition (for $n = 0$ this reduces to the $\text{TD}(0)$ case of only considering the current transition).

```

1 # parameters
2 num_episodes = 3000
3 alpha = 0.1
4 gamma = 0.9
5 epsilon = 0.5
6 on_policy = True # SARSA or Q-Learning
7 n = 2           # length of trace to use
8
9 # value initialisation
10 Q = np.random.uniform(0, 1e-5, [env.observation_space.n, env.action_space.n]) # noisy
11 Q = np.zeros([env.observation_space.n, env.action_space.n])                  # neutral
12 V = np.zeros([env.observation_space.n])
13
14 if on_policy:
15     # policies for SARSA
16     # vvvvvvvvvvvvvvvvvvvvvvvv
17     sample_policy = QPolicy(Q, epsilon)
18     learned_policy = sample_policy
19     plot_labels.append(f"SARSA (n={n}, alpha={alpha}, epsilon={epsilon})")
20     # ~~~~~
21 else:
22     # policies for Q-Learning
23     # vvvvvvvvvvvvvvvvvvvvvvvv
24     sample_policy = QPolicy(Q, epsilon)
25     td_epsilon = 0.01
26     learned_policy = QPolicy(Q, td_epsilon)
27     plot_labels.append(f"Q-Learning (n={n}, alpha={alpha}, epsilon={epsilon}|{td_epsilon})")
28     # ~~~~~
29
30 for episode in range(num_episodes):
31     state, _ = env.reset()
32     reward_sum = 0
33     done_n = 0
34
35     # trace of the last n + 1 transitions (state, action, reward, next_action)
36     trace = np.zeros((n + 1, 4), dtype=int)
37
38     # learning a policy
39     for t in itertools.count():
40         action = sample_policy.sample(state)
41         next_state, reward, term, trun, _ = env.step(action)
42         done = term or trun

```


SARSA(n)

Exercise

- Implement SARSA(n) by filling in the `n_step_return`, `v_target` and `q_target` in the skeleton.
- Run this on the 8x8 environment and compare different values of n in terms of performance (how quickly the agent learns) and run time (roughly how fast/slow everything is running). What effects do you observe?

Q-Learning

For Q-Learning there is a problem because we want to evaluate a different policy (π^*) than the one used for sampling (π). That means that the n -step return was sampled with the “wrong” policy (π^* might never take some of the actions sampled from π) and so is not representative for π^* . This bias can be corrected by adding an *importance sampling* factor (a general technique in Monte Carlo methods to correct for sampling from a “wrong” distribution) to the learning rate α

$$\rho_t = \frac{\pi^*(a_t | s_t) \pi^*(a_{t+1} | s_{t+1}) \dots \pi^*(a_{t+n} | s_{t+n})}{\pi(a_t | s_t) \pi(a_{t+1} | s_{t+1}) \dots \pi(a_{t+n} | s_{t+n})} = \frac{\prod_{k=0}^n \pi^*(a_{t+k} | s_{t+k})}{\prod_{k=0}^n \pi(a_{t+k} | s_{t+k})}$$
$$\alpha_t = \alpha \rho_t .$$

Exercise

- Implement TD(n) Q-Learning by
 - using different `sample_policy` and `learned_policy`
 - computing the importance sampling factor ρ (`policy[s][a]` is giving you the probability of taking action `a` in state `s`).
 - modifying the updates accordingly.
 - Test TD(n) Q-Learning with different values for n and exploration in the sampling policy on the 8x8 environment. *Hint:* Use some small non-zero value for `epsilon` in the learned policy to make sure the importance sampling factors are not (almost) all zero. What do you observe?
-