COMP3667: Reinforcement learning practical 6

Temporal-Difference Learning

robert.lieck@durham.ac.uk

1 Overview

Welcome to the sixth reinforcement learning practical. In this practical, we will be experimenting with different TD methods to better understand their characteristics, advantages, and drawbacks. In particular, we will

- evaluate on-policy SARSA(0) and off-policy Q-Learning

- experiment with n-step TD learning

- compare the performance of these methods in different environments

- understand when/how/why they may fail to work and how to fix them.

2 Setup

For this practical, you will need a basic Python environment with `numpy`, `matplotlib`, and OpenAI `gym` (version 0.20.0). You will not need a GPU. You can find a colab notebook with the complete setup code **here**.

</> copy code

```
1  !pip install 'gym[box2d]==0.20.0'
2  import gym
3  import numpy as np
4  import matplotlib.pyplot as plt
5  import matplotlib.font_manager
6  import itertools
7  from IPython import display
```

As before, we will use different versions of the frozen lake gym environment:

</> copy code

```
1   name = 'FrozenLake-v1'
2   env = gym.make(name, is_slippery=False)                 # 4x4
3   env = gym.make(name, map_name="8x8", is_slippery=False) # 8x8
4   env = gym.make(name, desc=["SFHH",
5                              "HFFH",
6                              "HHFF",
7                              "HHHG",], is_slippery=False) # custom
8   env.seed(742)
9   env.action_space.seed(742)
10  LEFT, DOWN, RIGHT, UP = 0, 1, 2, 3
```

You can use this function (slightly modified from previous practicals) for plotting the environment

</> copy code

```
1  def plot(env, v=None, policy=None, col_ramp=1, dpi=175, draw_vals=False, mark_ice=True):
2      # set up plot
3      plt.rcParams['figure.dpi'] = dpi
4      plt.rcParams.update({'axes.edgecolor': (0.32,0.36,0.38)})
5      plt.rcParams.update({'font.size': 4 if env.env.nrow == 8 else 7})
6      gray = np.array((0.32,0.36,0.38))
7      plt.figure(figsize=(3, 3))
8      ax = plt.gca()
```

```
9      ax.set_xticks(np.arange(env.env.ncol)-.5)
10     ax.set_yticks(np.arange(env.env.nrow)-.5)
11     ax.set_xticklabels([])
12     ax.set_yticklabels([])
13     plt.grid(color=(0.42,0.46,0.48), linestyle=':')
14     ax.set_axisbelow(True)
15     ax.tick_params(color=(0.42,0.46,0.48),
16                    which='both', top='off', left='off', right='off', bottom='off')
17     # use zero value as dummy if not provided
18     if v is None:
19         v = np.zeros(env.nS)
20     # plot values
21     plt.imshow(1-v.reshape(env.env.nrow,env.env.ncol)**col_ramp,
22                cmap='gray', interpolation='none',
23                clim=(0,1), zorder=-1)
24     # go through states
25     for s in range(env.nS):
26         x, y = s % env.env.nrow, s // env.env.ncol
27         # print numeric values
28         if draw_vals and v[s] > 0:
29             vstr = '{0:.1e}'.format(v[s]) if env.env.nrow == 8 else '{0:.6f}'.format(v[s])
30             plt.text(x - 0.45, y + 0.45, vstr, color=(0.2, 0.8, 0.2), fontname='Sans')
31         # mark ice, start, goal
32         if env.desc.tolist()[y][x] == b'F':
33             plt.text(x-0.45,y-0.3, 'ice', color=(0.5, 0.6, 1), fontname='Sans')
34             if mark_ice:
35                 ax.add_patch(plt.Circle((x, y), 0.2, color=(0.7, 0.8, 1), zorder=0))
36         elif env.desc.tolist()[y][x] == b'S':
37             plt.text(x-0.45,y-0.3, 'start',color=(0.2,0.5,0.5), fontname='Sans',
38                      weight='bold')
39         elif env.desc.tolist()[y][x] == b'G':
40             plt.text(x-0.45,y-0.3, 'goal', color=(0.7,0.2,0.2), fontname='Sans',
41                      weight='bold')
42             continue # don't plot policy for goal state
43         else:
44             continue # don't plot policy for holes
45         # plot policy
46         def plot_arrow(x, y, dx, dy, v, scale=0.4):
47             plt.arrow(x, y, scale * float(dx), scale * float(dy), color=gray+0.2*(1-v),
48                       head_width=0.1, head_length=0.1, zorder=1)
49         if policy is not None:
50             if policy[s, 0] > 0.0: plot_arrow(x, y, -policy[s, 0],        0., v[s]) # left
51             if policy[s, 1] > 0.0: plot_arrow(x, y,        0.,  policy[s, 1], v[s]) # down
52             if policy[s, 2] > 0.0: plot_arrow(x, y,  policy[s, 2],        0., v[s]) # right
53             if policy[s, 3] > 0.0: plot_arrow(x, y,        0., -policy[s, 3], v[s]) # up
54     plt.show()
```

and these two helper classes to define hard-coded policies or policies using Q-values

</> copy code

```
1   class QPolicy:
2       def __init__(self, Q, epsilon, values=False):
3           self.Q = Q
4           self.epsilon = epsilon
5           self.values = values
6       def sample(self, state):
7           if np.random.rand() > self.epsilon:
8               best_actions = np.argwhere(self.Q[state]==np.max(self.Q[state])).flatten()
9               return np.random.choice(best_actions)
10          else:
11              return env.action_space.sample()
```

```python
    def __getitem__(self, item):
        state, action = item
        if self.values:
            return self.Q[state, action] / (self.Q[state].sum() + 1e-10)
        else:
            best_actions = np.argwhere(self.Q[state]==np.max(self.Q[state])).flatten()
            p = int(action in best_actions) / len(best_actions)
            return (1 - self.epsilon) * p + self.epsilon / len(self.Q[state])
class HardCodedPolicy:
    def __init__(self, state_action_map):
        self.state_action_map = state_action_map
    def sample(self, state):
        if state in self.state_action_map:
            return np.random.choice(self.state_action_map[state])
        else:
            return np.random.choice(4)
    def __getitem__(self, item):
        state, action = item
        if state in self.state_action_map:
            if action in self.state_action_map[state]:
                return 1 / len(self.state_action_map[state])
            else:
                return 0
        else:
            return 1 / 4
```

Finally, these functions can be used to plot the learning progress

</> copy code

```python
# (using global variables in functions)
def update_plot(mod):
    reward_list[experiment_id].append(reward_sum)
    aoc[experiment_id] += reward_sum
    test_reward_list[experiment_id].append(test_reward_sum)
    test_aoc[experiment_id] += test_reward_sum
    if episode % mod == 0:
        plot_data[experiment_id].append([episode,
                                         np.array(reward_list[experiment_id]).mean(),
                                         np.array(test_reward_list[experiment_id]).mean()])
        reward_list[experiment_id] = []
        test_reward_list[experiment_id] = []
        for i in range(len(plot_data)):
            color=next(plt.gca()._get_lines.prop_cycler)['color']
            plt.plot([x[0] for x in plot_data[i]],
                     [x[1] for x in plot_data[i]],
                     '-', color=color,
                     label=f"{plot_labels[i]}, AOC: {aoc[i]}|{test_aoc[i]}")
            plt.plot([x[0] for x in plot_data[i]],
                     [x[2] for x in plot_data[i]], '--', color=color)
        plt.xlabel('Episode number')
        plt.ylabel('Episode reward')
        plt.legend()
        display.clear_output(wait=True)
        plt.show()
def next_experiment():
    reward_list.append([])
    aoc.append(0)
    test_reward_list.append([])
    test_aoc.append(0)
    plot_data.append([])
    return experiment_id + 1
```

using the following global variables (re-evaluate this cell to reset them)

```
1  reward_list = [[]]
2  aoc = [0]
3  test_reward_list = [[]]
4  test_aoc = [0]
5  plot_data = [[]]
6  plot_labels = []
7  experiment_id = 0
```

3 On-policy and off-policy learning with TD(0)

Recap TD Learning

Remember our 0-step temporal difference (TD) targets from the lecture, which can be computed for any (also partial) episodes

$$
V_\pi(s_t) = \sum_{a_t \in \mathcal{A}} \pi(a_t \mid s_t) \sum_{s_{t+1} \in \mathcal{S}} p(s_{t+1} \mid s_t, a_t) \Big[\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, V_\pi(s_{t+1}) \Big]
$$

$$
= \mathbb{E}_{a_t \sim \pi(a_t \mid s_t)} \mathbb{E}_{s_{t+1} \sim p(s_{t+1} \mid s_t, a_t)} \Big[\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, V_\pi(s_{t+1}) \Big]
$$

$$
\approx \underbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, V_\pi(s_{t+1})}_{\text{TD(0) target}}
$$

$$
Q_{\pi^*}(s_t, a_t) = \sum_{s_{t+1} \in \mathcal{S}} p(s_{t+1} \mid s_t, a_t) \Big[\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma \sum_{a_{t+1} \in \mathcal{A}} \pi^*(a_{t+1} \mid s_{t+1})\, Q_{\pi^*}(s_{t+1}, a_{t+1}) \Big]
$$

$$
= \mathbb{E}_{s_{t+1} \sim p(s_{t+1} \mid s_t, a_t)} \Big[\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, \mathbb{E}_{a_{t+1} \sim \pi^*(a_{t+1} \mid s_{t+1})} Q_{\pi^*}(s_{t+1}, a_{t+1}) \Big]
$$

$$
\approx \underbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, Q_{\pi^*}(s_{t+1}, a_{t+1})}_{\text{TD(0) target}} \ .
$$

At a particular time t we are in state s_t and take action $a_t \sim \pi(a_t \mid s_t)$ sampled from the policy π. We then end up in state $s_{t+1} \sim p(s_{t+1} \mid s_t, a_t)$ based on the environment's transition function. If we are interested in learning state-action values (i.e. solving the control problem), we additionally need to consider the following action $a_{t+1} \sim \pi^*(a_{t+1} \mid s_{t+1})$ based on the policy π^* that we want to learn or evaluate.

The TD(0) targets are noisy "snapshots" of how the values should look like based on the current transition at time t. The difference between the TD target and our current value estimate gives us a noisy TD error signal that tells us "how far off" our estimates are. This can be use to update our value estimates with a learning rate α (similar to SGD) to improve them

$$
V_\pi(s_t) \leftarrow V_\pi(s_t) + \alpha \Big[\underbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, V_\pi(s_{t+1})}_{\text{TD(0) target}} - V_\pi(s_t) \Big]
$$

$$
Q_{\pi^*}(s_t, a_t) \leftarrow Q_{\pi^*}(s_t, a_t) + \alpha \Big[\underbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma\, Q_{\pi^*}(s_{t+1}, a_{t+1})}_{\text{TD(0) target}} - Q_{\pi^*}(s_t, a_t) \Big] \ .
$$

On-policy SARSA(0)

We can evaluate and improve our policy "on the go". This is called *on-policy* learning and it means that the policy π we use for sampling is the same as the policy π^* we are evaluating and learning.

Exercise

- Implement SARSA(0) by filling in the TD(0) targets and updates in the skeleton below. (The rest of the skeleton is for collecting episodes, evaluating the learned policy and plotting everything. The dashed line is the reward for the learned policy, the solid line is for the sampling policy.)

- Run a couple of evaluations with different values for the learning rate `alpha` and the exploration `epsilon` on the 4x4 environment. Use both "noisy" and "neutral" initialisations for Q (commenting in/out the respective lines in the code). What effects do you observe? With what parameters does the agent learn best?

</> copy code

```
1   # parameters
2   num_episodes = 3000
3   alpha = 0.1
4   gamma = 0.9
5   epsilon = 0.5
6
7   Q = np.random.uniform(0, 1e-5, [env.observation_space.n, env.action_space.n]) # noisy
8   Q = np.zeros([env.observation_space.n, env.action_space.n])                    # neutral
9   V = np.zeros([env.observation_space.n])
10
11  # policies
12  sample_policy = QPolicy(Q, epsilon)
13  learned_policy = sample_policy
14  plot_labels.append(f"SARSA (alpha={alpha}, epsilon={epsilon})")
15
16  for episode in range(num_episodes):
17      state = env.reset()
18      reward_sum = 0
19      # learning a policy
20      for t in itertools.count():
21          action = sample_policy.sample(state)
22          next_state, reward, done, _ = env.step(action)
23          next_action = learned_policy.sample(next_state)
24          # TD(0) targets
25          v_target = ...                    # FILL IN HERE!
26          q_target = ...                    # FILL IN HERE!
27          # updates
28          s, a = state, action
29          V[s] += ...                       # FILL IN HERE!
30          Q[s, a] += ...                    # FILL IN HERE!
31
32          reward_sum += reward
33          if done:
34              break
35          state = next_state
36
37      # testing the learned policy
38      state = env.reset()
39      test_reward_sum = 0
40      while True:
41          action = learned_policy.sample(state)
42          next_state, reward, done, _ = env.step(action)
43          test_reward_sum += reward
44          state = next_state
45          if done:
46              break
47
48      update_plot(int(np.ceil(num_episodes / 20)))
49
50  env.close()
51  experiment_id = next_experiment()
52  print("Sampling policy and values")
53  plot(env, v=V, policy=sample_policy, draw_vals=True)
54  print("Learned policy and optimal/max values")
```

```
55    plot(env, v=Q.max(axis=1), policy=learned_policy, draw_vals=True)
```

Off-policy Q-Learning

Sometimes, we would like to generate samples with one policy π (typically an exploratory policy) but then use the samples to learn another policy π^* (typically a near-optimal policy). This is possible with *off-policy* methods, such as Q-Learning.

3.0.1 Exercise

- Implement Q-Learning by using a different policy for the `learned_policy` (with lower `epsilon`) then for the `sample_policy`.

  - *Note:* Strictly speaking, the `sample_policy` does not change in Q-Learning. You can achieve this by using `epsilon=1` in the `sample_policy` to select actions randomly. For other values of `epsilon`, the `sample_policy` "peeks" into the values of the learned policy, which lets it profit from that learning (but is not a clean implementation).

  - *Bonus exercise:* Learn the `sample_policy` using SARSA(0) while learning the `learned_policy` with Q-Learning (this requires maintaining two copies of values estimates, performing separate updates for both etc).

- Run the evaluation several times on the 4x4 and 8x8 environment. Try to get the fastest and most reliable learning by tweaking `epsilon`. How does the performance (episode reward) of the sampling policy (solid line) compare to that of the learned policy (dashed lines)? How is this different from SARSA(0)?

Expected TD(0) targets

The normal TD targets to estimate state-action values are computed by *sampling* from the policy π^* that is to be learned. This sampling step increases the noise in the TD error signal (in addition to the noise we already have due to sampling episodes). Instead, one can use the *expected* TD targets

$$\underbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma \sum_{a_{t+1} \in \mathcal{A}} \pi^*(a_{t+1} \mid s_{t+1}) \, Q_{\pi^*}(s_{t+1}, a_{t+1})}_{\text{expected TD(0) target}} \ .$$

Exercise

- In the `q_target` replace the sampled value with an expectation over possible actions the `learned_policy` could take. *Note:* `learned_policy[next_state, a]` is giving you the probability of taking action `a` in state `next_state`.

- Compare the performance when using the expected TD(0) target to using the sampled one.

4 TD(n)

In TD(0) we do not look ahead and receiving a reward does only affect the value estimate of the current state and action. All the future expected rewards are approximated by using the current value estimates. However, we can improve on that by instead look n steps ahead (in hindsight) and taking the n next steps into account for computing TD targets. These TD(n) targets are

$$V_\pi(s_t) \approx \underbrace{\overbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma \, \mathcal{R}^{a_{t+1}}_{s_{t+1} s_{t+2}} + \ldots + \gamma^n \, \mathcal{R}^{a_{t+n}}_{s_{t+n} s_{t+n+1}}}^{n\text{-step return}} + \gamma^{n+1} V_\pi(s_{t+n+1})}_{\text{TD}(n) \text{ target}}$$

$$Q_{\pi^*}(s_t, a_t) \approx \underbrace{\mathcal{R}^{a_t}_{s_t s_{t+1}} + \gamma \, \mathcal{R}^{a_{t+1}}_{s_{t+1} s_{t+2}} + \ldots + \gamma^n \, \mathcal{R}^{a_{t+n}}_{s_{t+n} s_{t+n+1}} + \gamma^{n+1} Q_{\pi^*}(s_{t+n+1}, a_{t+n+1})}_{\text{TD}(n) \text{ target}} \ .$$

To compute them in hindsight, we have to store a trace of the last $n + 1$ transitions (including the current transition that is also used in TD(0)). The modified skeleton below is providing a trace of length $n+1$ containing the last n steps and the current transition (for $n = 0$ this reduces to the TD(0) case of only considering the current transition).

```python
# parameters
num_episodes = 3000
alpha = 0.1
gamma = 0.9
epsilon = 0.5
n = 0

Q = np.random.uniform(0, 1e-5, [env.observation_space.n, env.action_space.n]) # noisy
Q = np.zeros([env.observation_space.n, env.action_space.n])                    # neutral
V = np.zeros([env.observation_space.n])

# policies
sample_policy = QPolicy(Q, epsilon)
learned_policy = sample_policy
plot_labels.append(f"SARSA (n={n}, alpha={alpha}, epsilon={epsilon})")

for episode in range(num_episodes):
    state = env.reset()
    reward_sum = 0
    trace = np.zeros((n + 1, 4), dtype=int)
    done_n = 0
    # learning a policy
    for t in itertools.count():
        action = sample_policy.sample(state)
        next_state, reward, done, _ = env.step(action)
        next_action = learned_policy.sample(next_state)
        trace[-1] = (state, action, reward, next_action)
        if t > n:
            # n-step targets
            n_step_return = ...  # FILL IN HERE!
            v_target = ...       # FILL IN HERE!
            q_target = ...       # FILL IN HERE!
            # updates
            s, a, _, _ = trace[0]
            V[s] += alpha * (v_target - V[s])
            Q[s, a] += alpha * (q_target - Q[s, a])

        reward_sum += reward
        done_n += done
        if done_n > n:
            break
        state = next_state
        trace = np.roll(trace, shift=-1, axis=0)

    # testing the learned policy
    state = env.reset()
    test_reward_sum = 0
    while True:
        action = learned_policy.sample(state)
        next_state, reward, done, _ = env.step(action)
        test_reward_sum += reward
        state = next_state
        if done:
            break

    update_plot(int(np.ceil(num_episodes / 20)))

env.close()
experiment_id = next_experiment()
```

```
60    print("Sampling policy and values")
61    plot(env, v=V, policy=sample_policy, draw_vals=True)
62    print("Learned policy and optimal/max values")
63    plot(env, v=Q.max(axis=1), policy=learned_policy, draw_vals=True)
```

SARSA(n)

Exercise

- Implement SARSA(n) by filling in the `n_step_return`, `v_target` and `q_target` in the skeleton.

- Run this on the 8x8 environment and compare different values of n in terms of performance (how quickly the agent learns) and run time (roughly how fast/slow everything is running). What effects do you observe?

Q-Learning

For Q-Learning there is a problem because we want to evaluate a different policy (π^*) than the one used for sampling (π). That means that the n-step return was sampled with the "wrong" policy (π^* might never take some of the actions sampled from π) and so is not representative for π^*. This bias can be corrected by adding an *importance sampling* factor (a general technique in Monte Carlo methods to correct for sampling from a "wrong" distribution) to the learning rate α

$$\rho_t = \frac{\pi^*(a_t \mid s_t)\,\pi^*(a_{t+1} \mid s_{t+1})\,\ldots\,\pi^*(a_{t+n} \mid s_{t+n})}{\pi(a_t \mid s_t)\,\pi(a_{t+1} \mid s_{t+1})\,\ldots\,\pi(a_{t+n} \mid s_{t+n})} = \frac{\prod_{k=0}^{n} \pi^*(a_{t+k} \mid s_{t+k})}{\prod_{k=0}^{n} \pi(a_{t+k} \mid s_{t+k})}$$

$$\alpha_t = \alpha \rho_t \;.$$

Exercise

- Implement TD(n) Q-Learning by

    - using different `sample_policy` and `learned_policy`
    - computing the importance sampling factor ρ (`policy[s, a]` is giving you the probability of taking action `a` in state `s`).
    - modifying the updates accordingly.

- Test TD(n) Q-Learning with different values for n and exploration in the sampling policy on the 8x8 environment. *Hint:* Use some small non-zero value for `epsilon` in the learned policy to make sure the importance sampling factors are not (almost) all zero. What do you observe?

———————————————————————————